Step 0: Create a new project

1. Open your browser and go to IBM Cloud Login: https://cloud.ibm.com/login
2. Sign in with your IBM Cloud credentials.
3. Expand Manage → click Access (IAM).

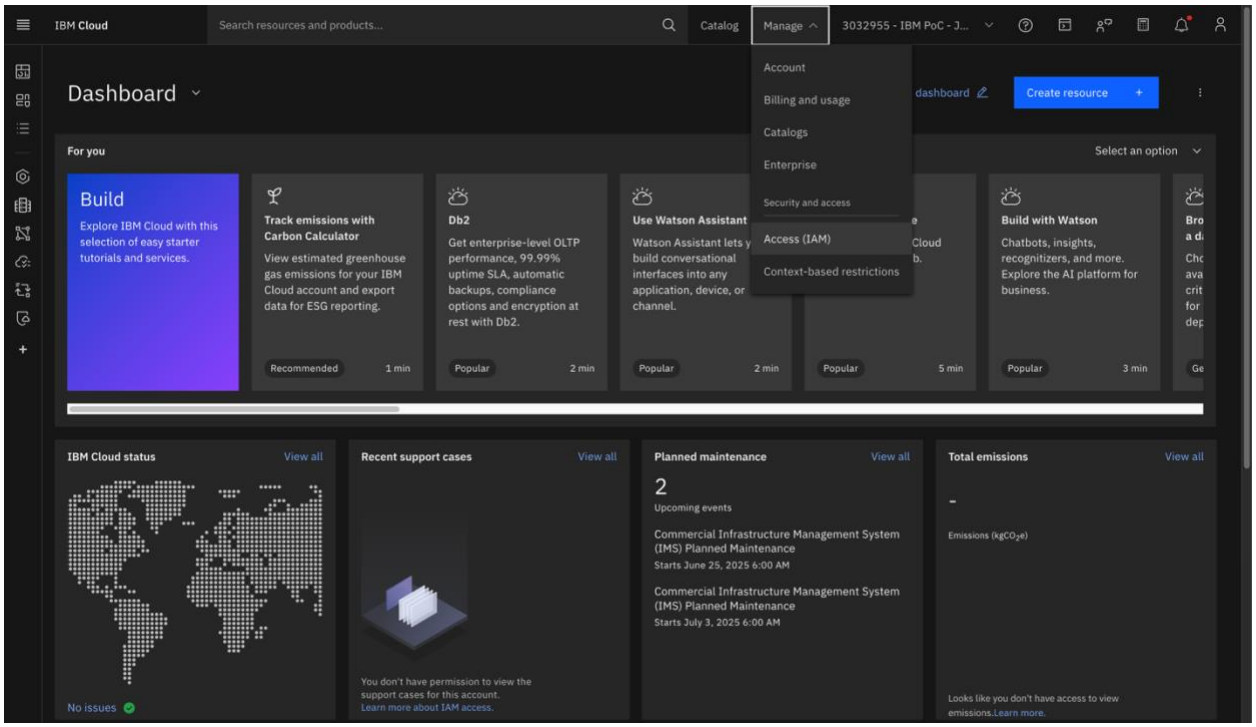

4. Select API Keys.

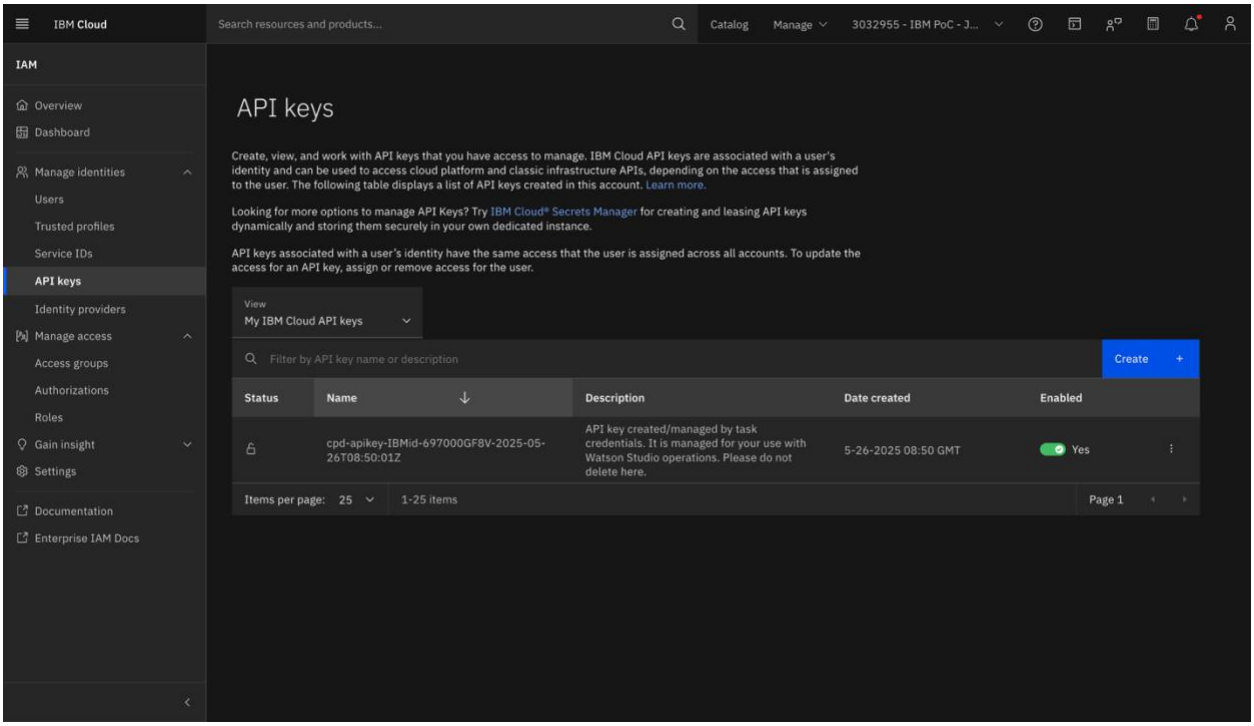

5. Click Create and enter the name of key: wx_apikey.

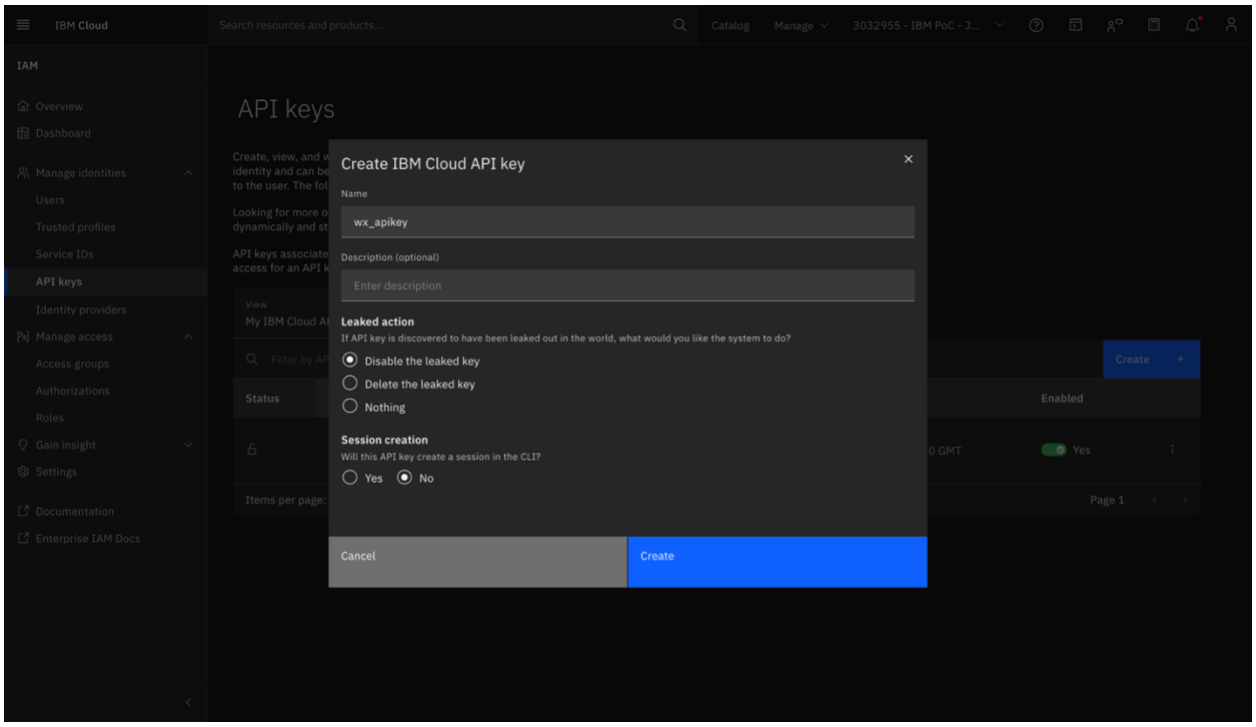

6. Copy the displayed API key and save it **securely**.

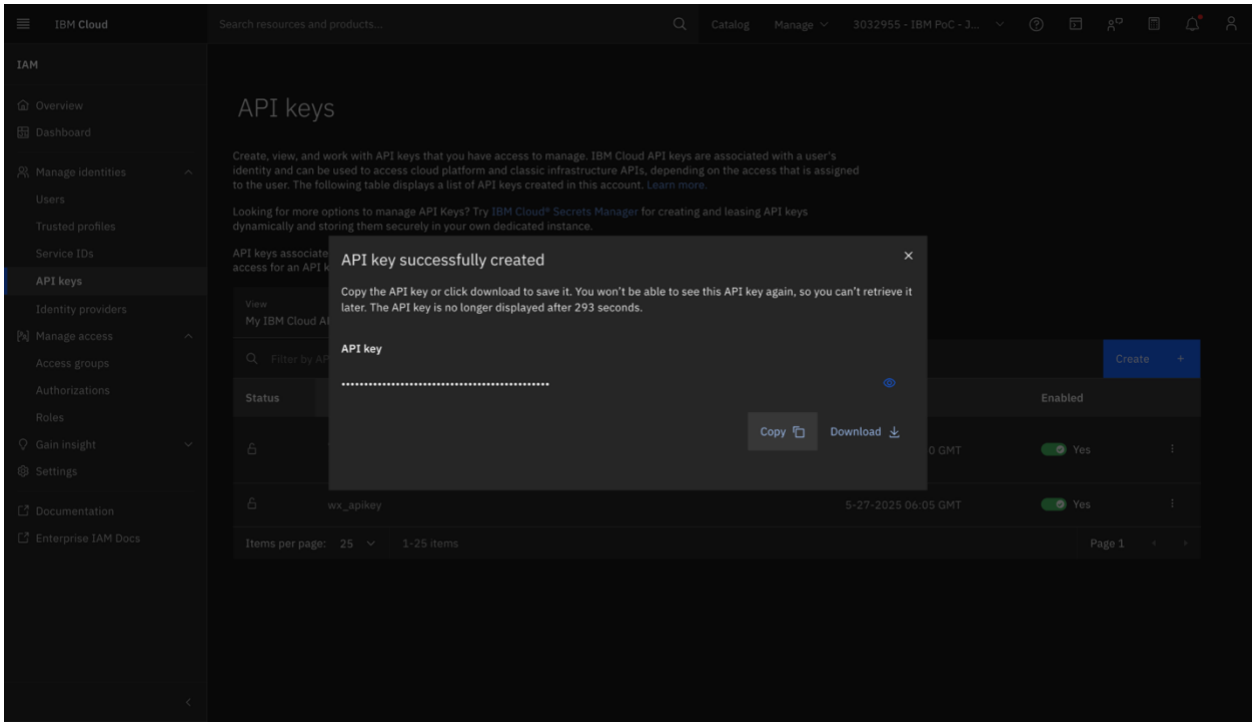

7. In the left-hand sidebar, click Resource list.

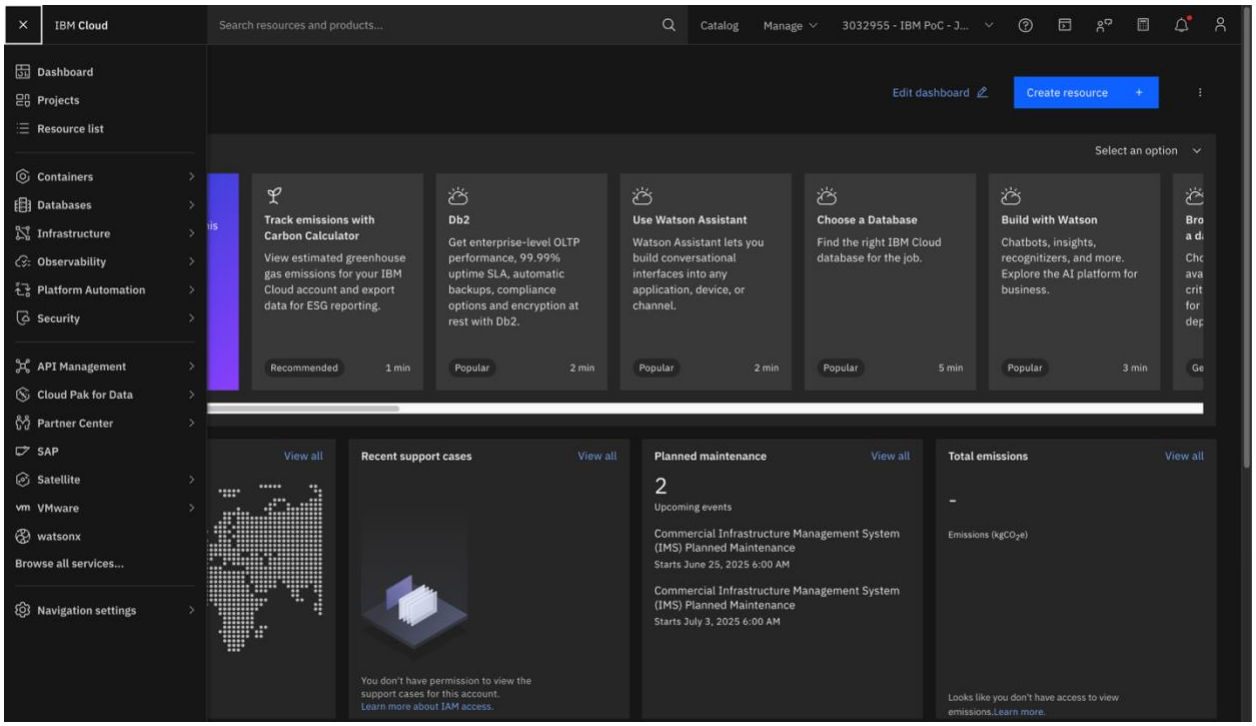

8. In the search box, type watsonx.ai, then from the results select your assigned service watsonx.ai group.

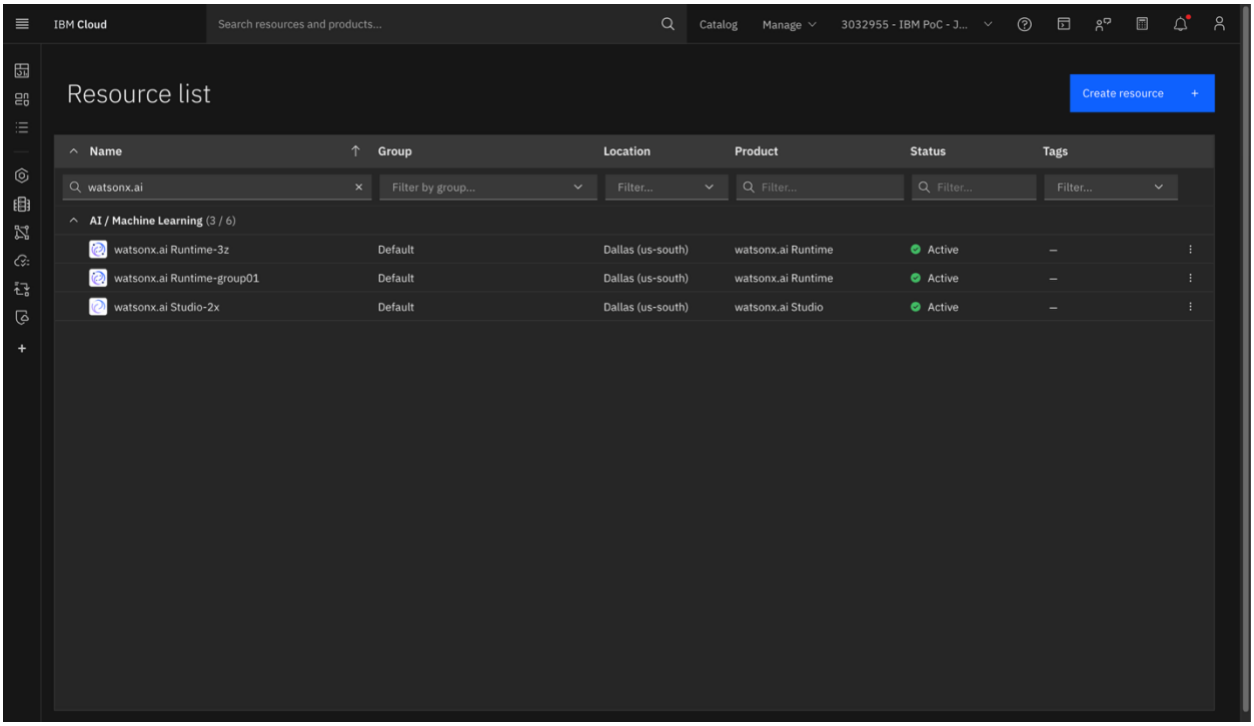

9. Click Launch in → **IBM watsonx** to open the watsonx console.

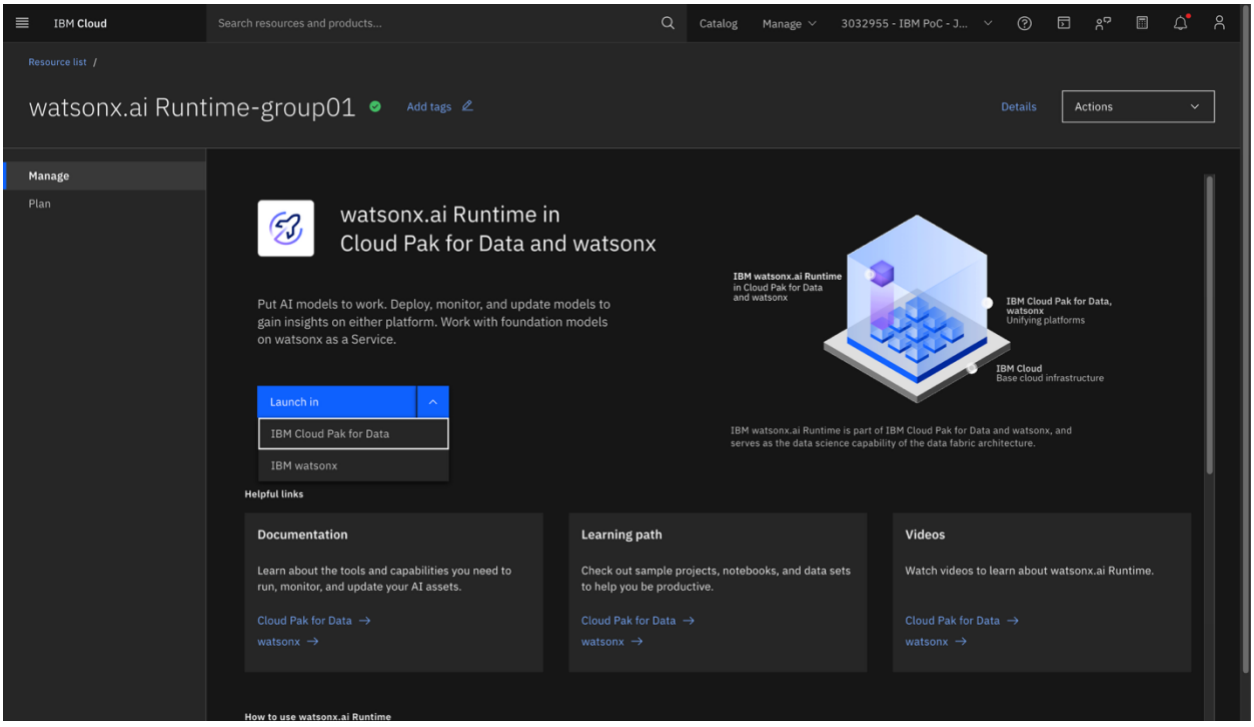

10. In the IBM watsonx console sidebar, click Projects.
11. Click the New project button in the top-right.
12. Ensure the New tab is selected.
13. In Name, enter <Your Name>'s Sandbox.
14. (Optional) Fill in Description and Tags as desired.

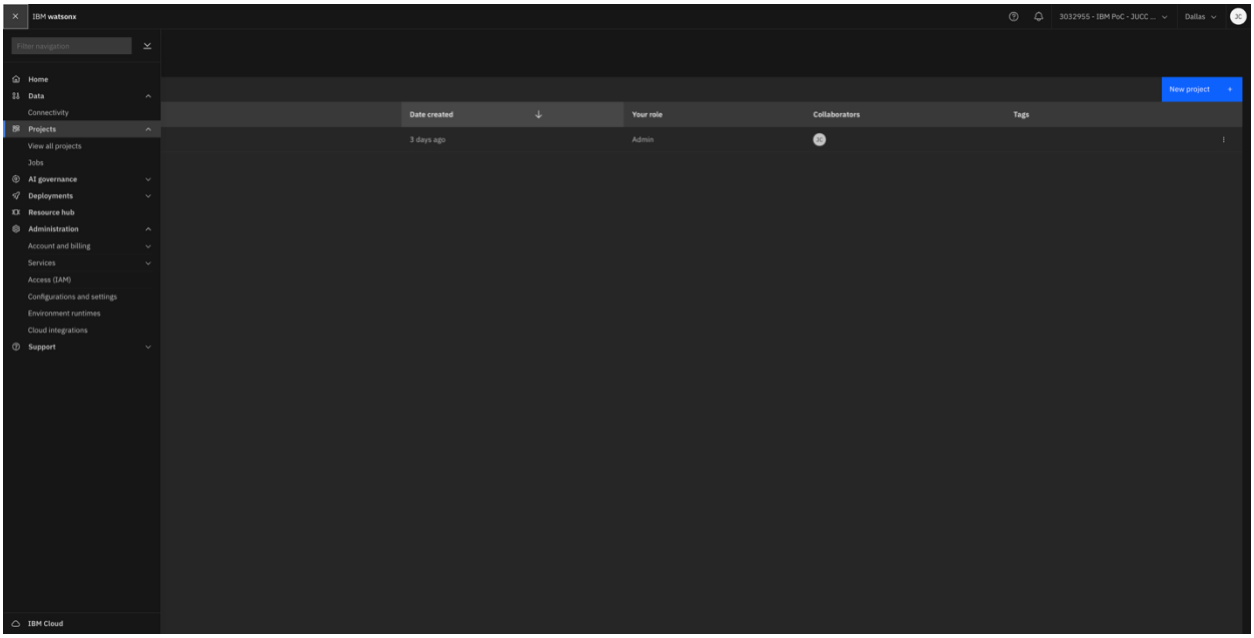

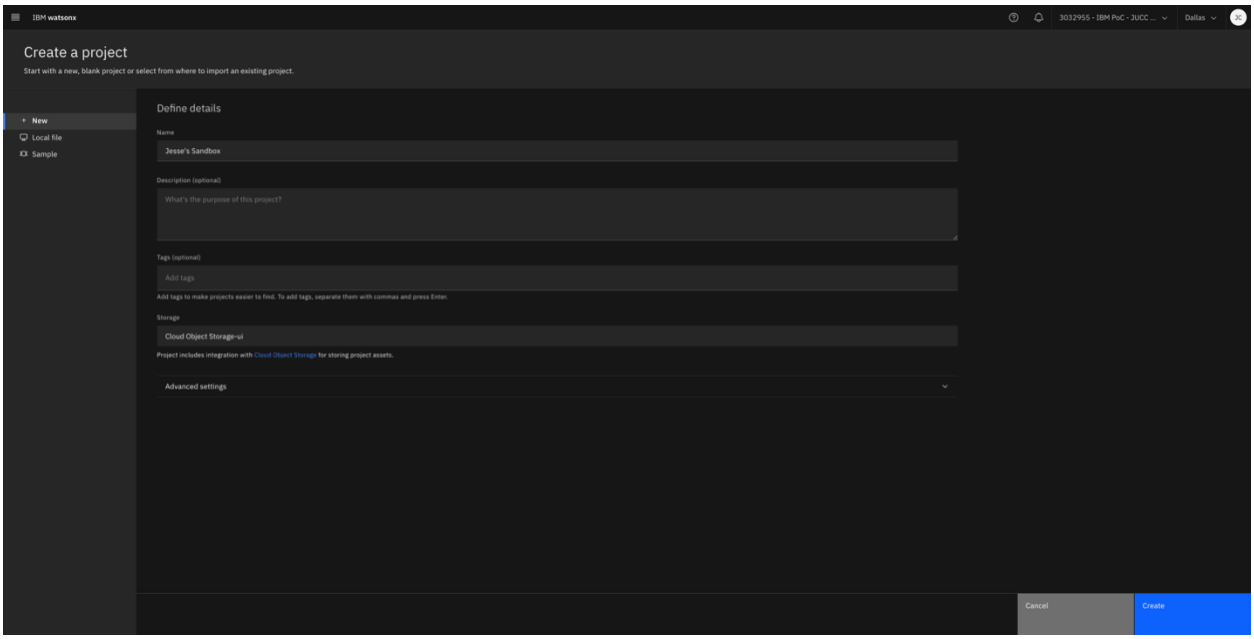

Step 0.5: Import Elasticsearch Connection

1. Click on the “New Asset” button in the project.

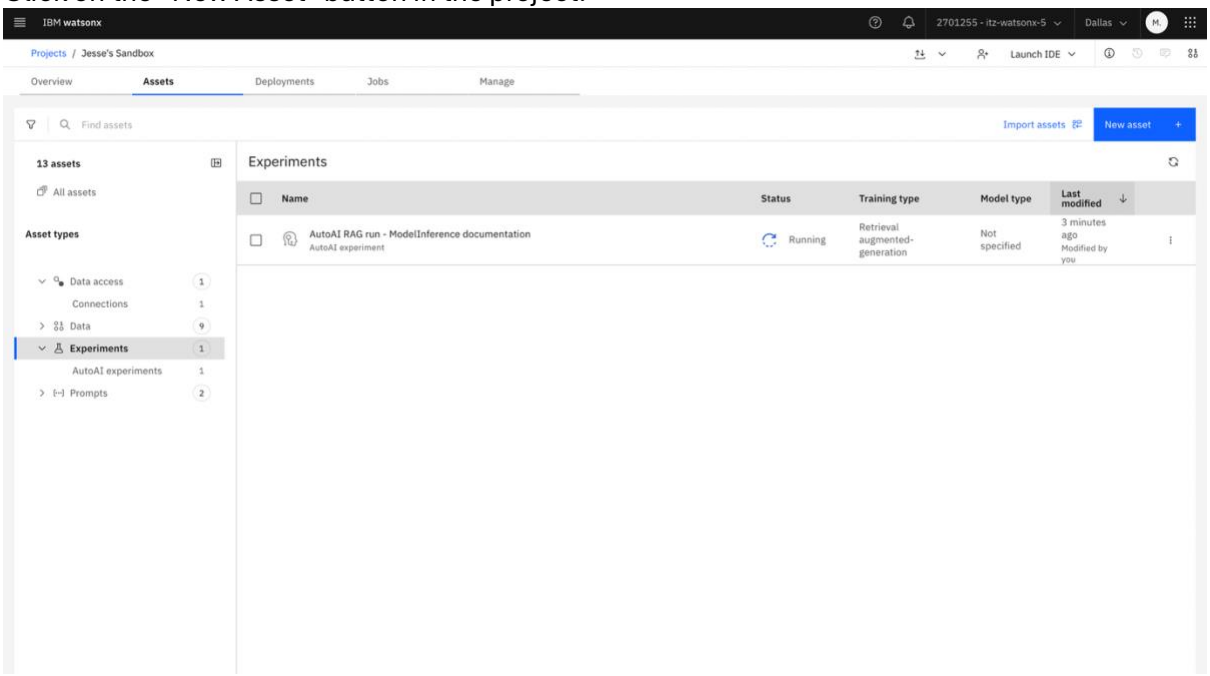

2. Now choose the option “Connect to a data source”.

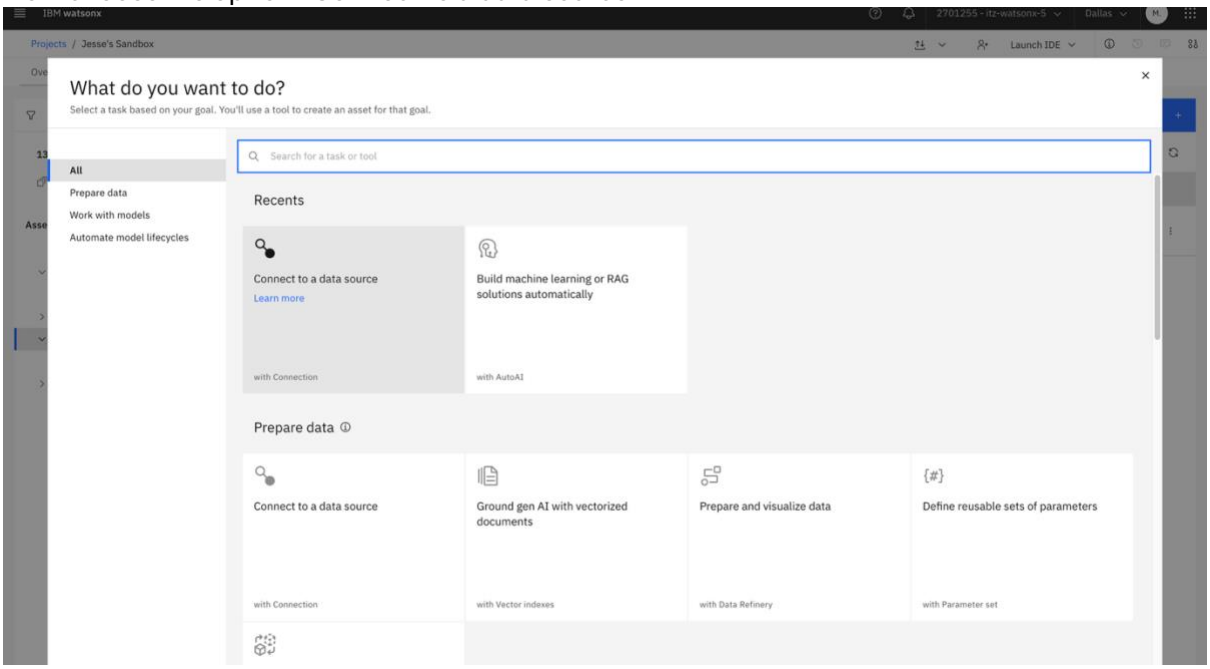

3. Select Elasticsearch from the options on left panel. Then click on Next.

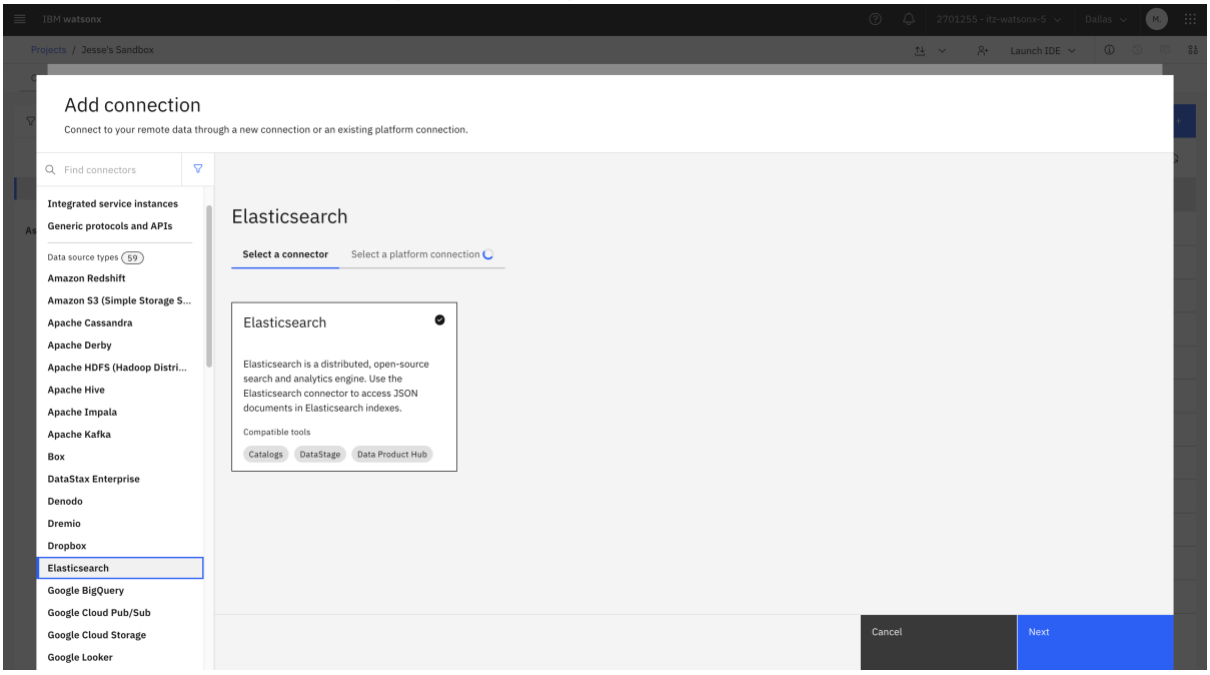

4. In a separate tab, navigate to the Databases for Elasticsearch service on IBM Cloud and copy the hostname, port, and TLS certificate under the Endpoints section of the Overview tab.

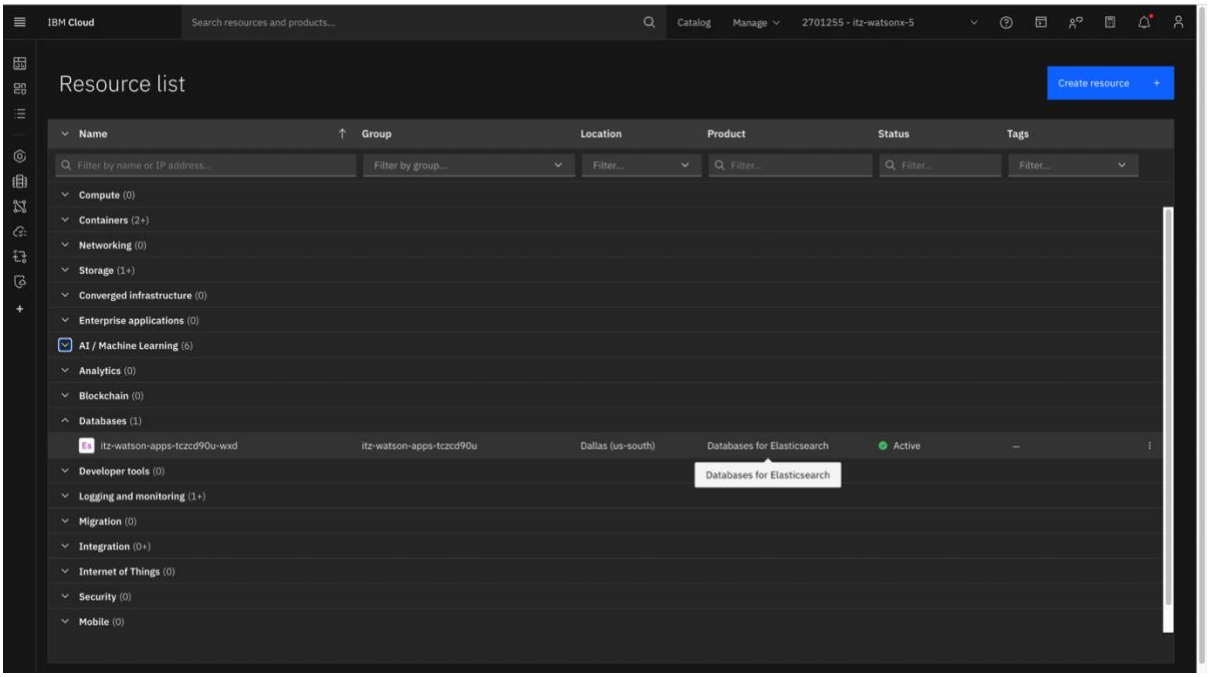

Copy the TLS certificate and add it to Elasticsearch connection on the watsonx.ai platform.

5. From the HTTPS tab you can note down the Hostname and Port number.

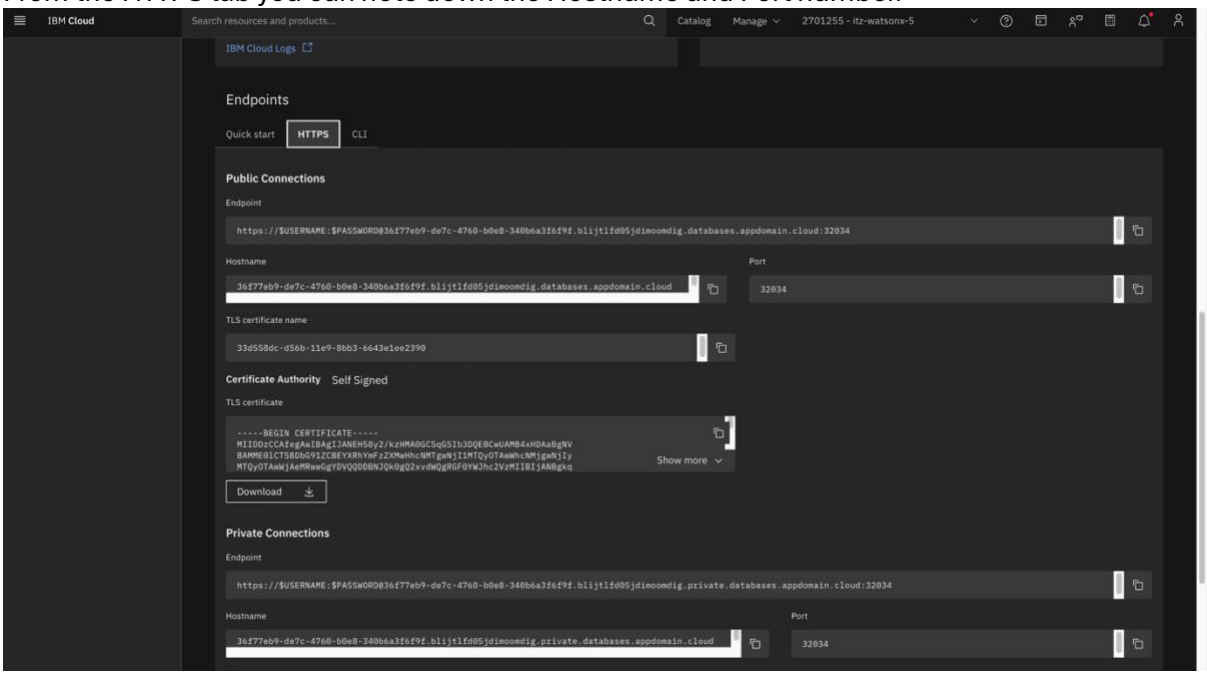

6. Go to the service credentials tab of the service, click on the arrow to expand it. Copy the username and password under **“connection.https.authentication”** in the service credential JSON.

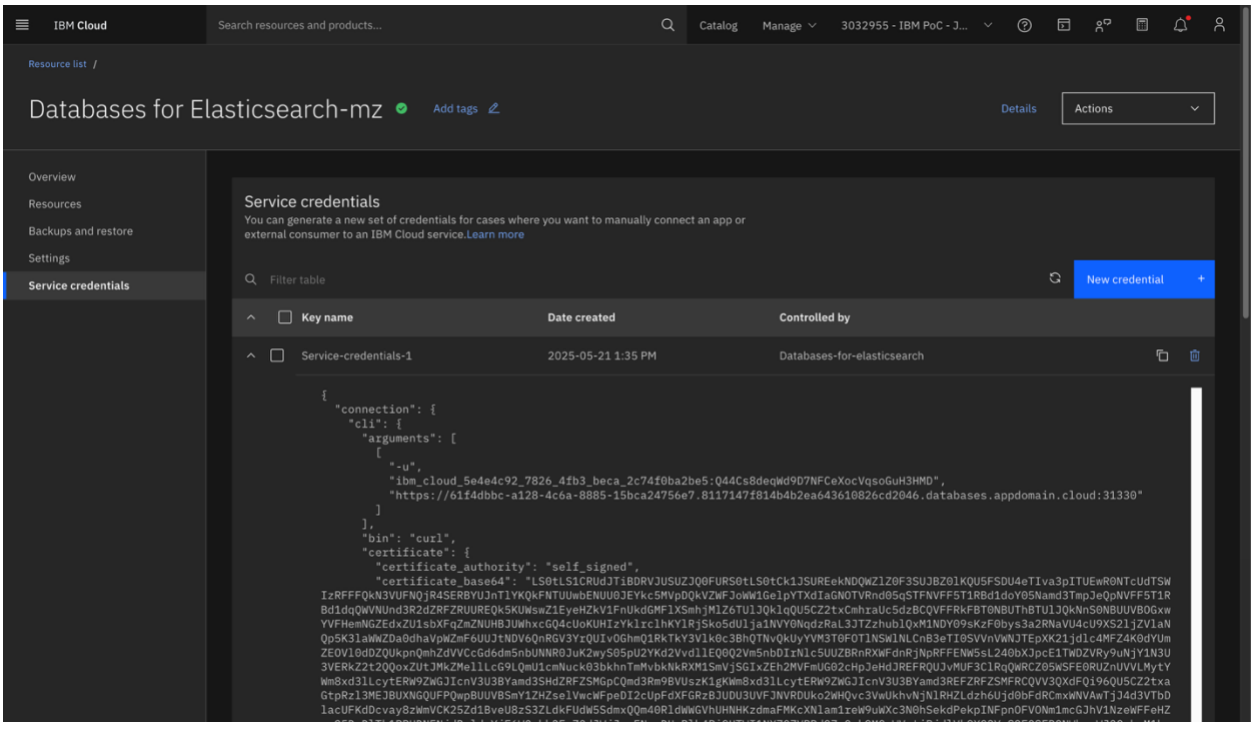

7. Return to the project connection and set the following fields. Give the connection a name of: “Elasticsearch connection”

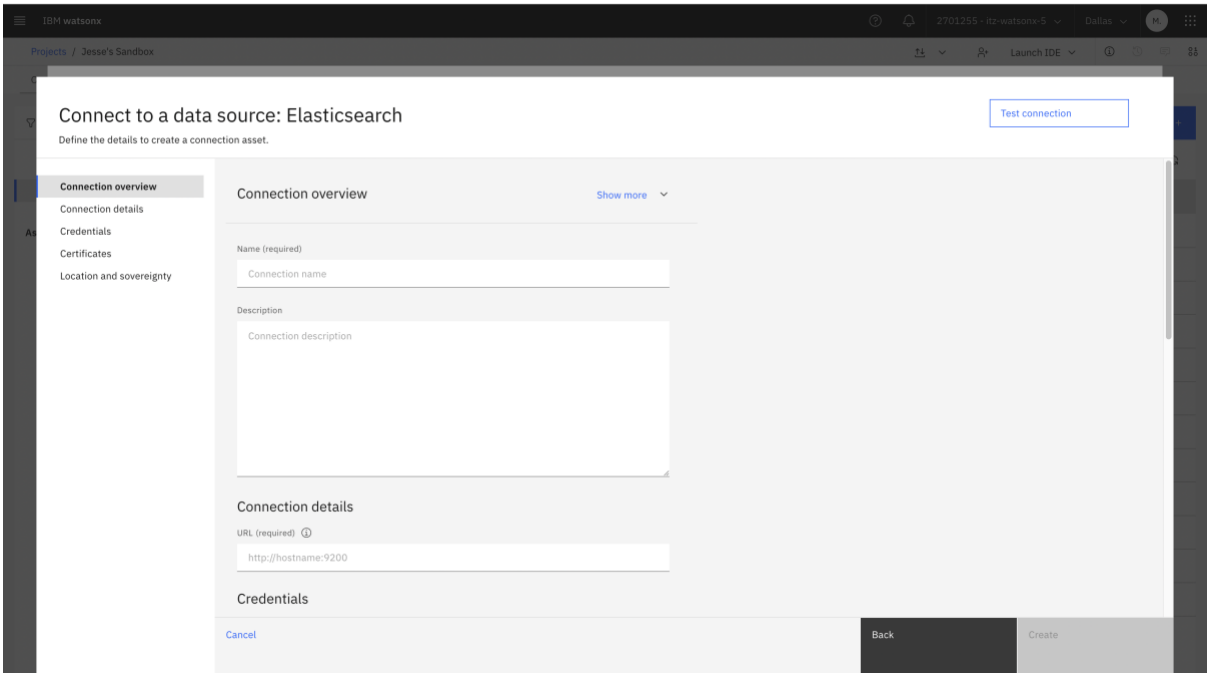

8. The URL should be formatted as https://{HOSTNAME}:{PORT}.

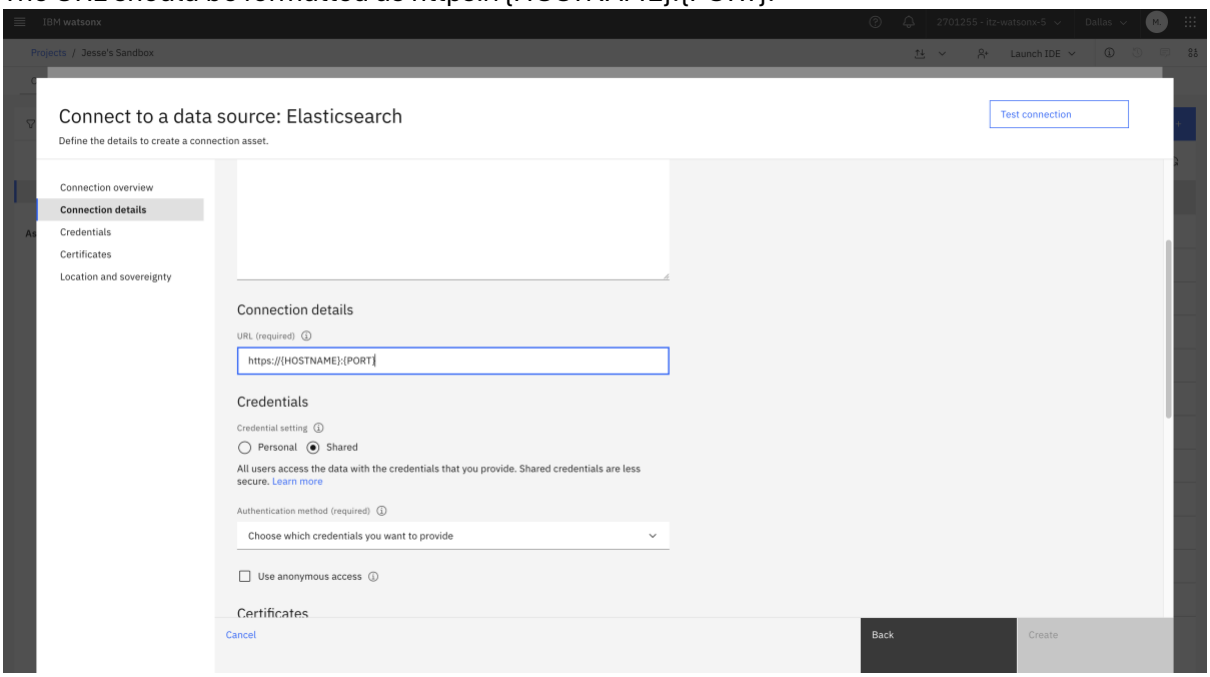

9. Now for authentication, from the dropdown select Username and password.

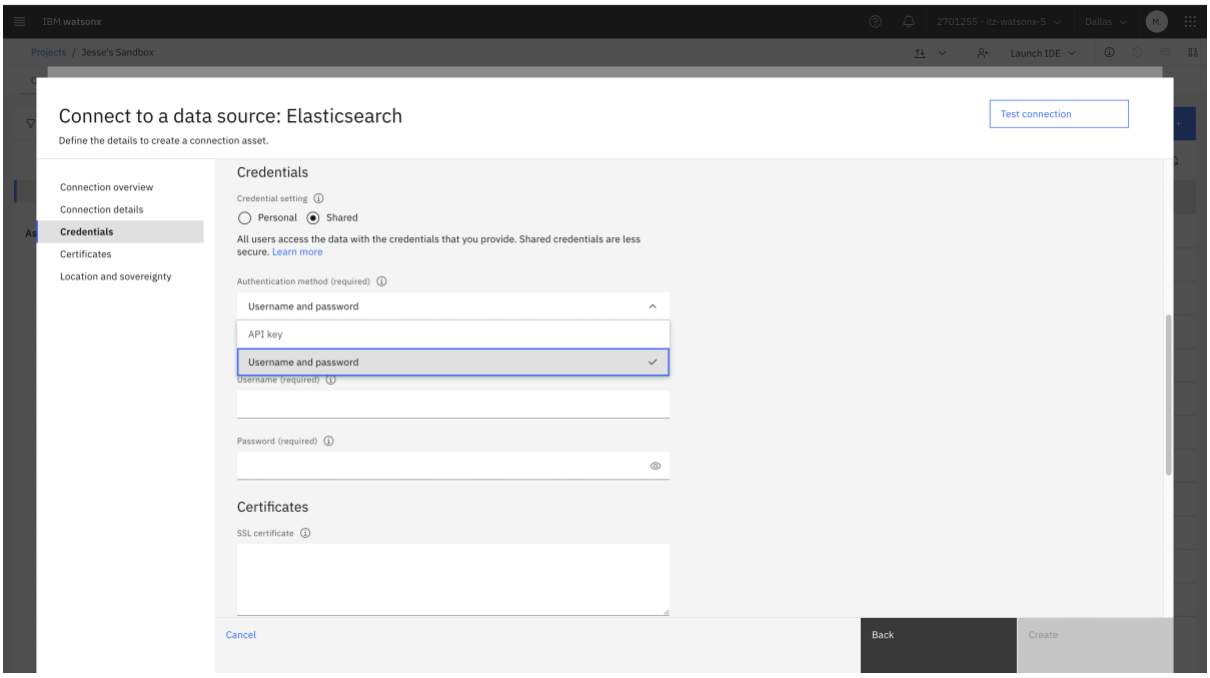

10. The username and password should be the ones copied from the service credentials.
11. The SSL certificate should be the TLS certificate.
12. Click on Test connection on the top, if that is successful then click on create.

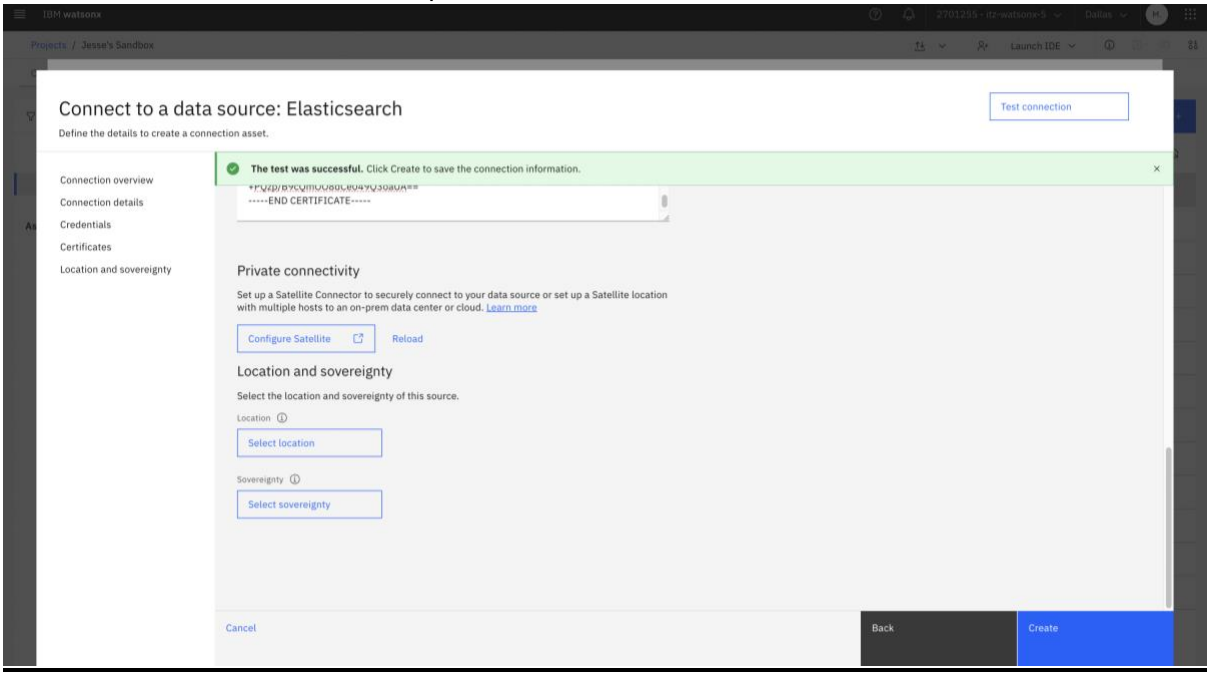

Step 0.5: Import Data Asset

1. Navigate to the GitHub repo at

   https://github.com/JesseChan5171/JUCC-workshop#

2. Click Code → Download ZIP (or run git clone)
3. In the extracted directory, locate the “evaluation_dataset” folder.
4. In your project, click the **Assets** tab.
5. Click **Import assets** in the top-right.

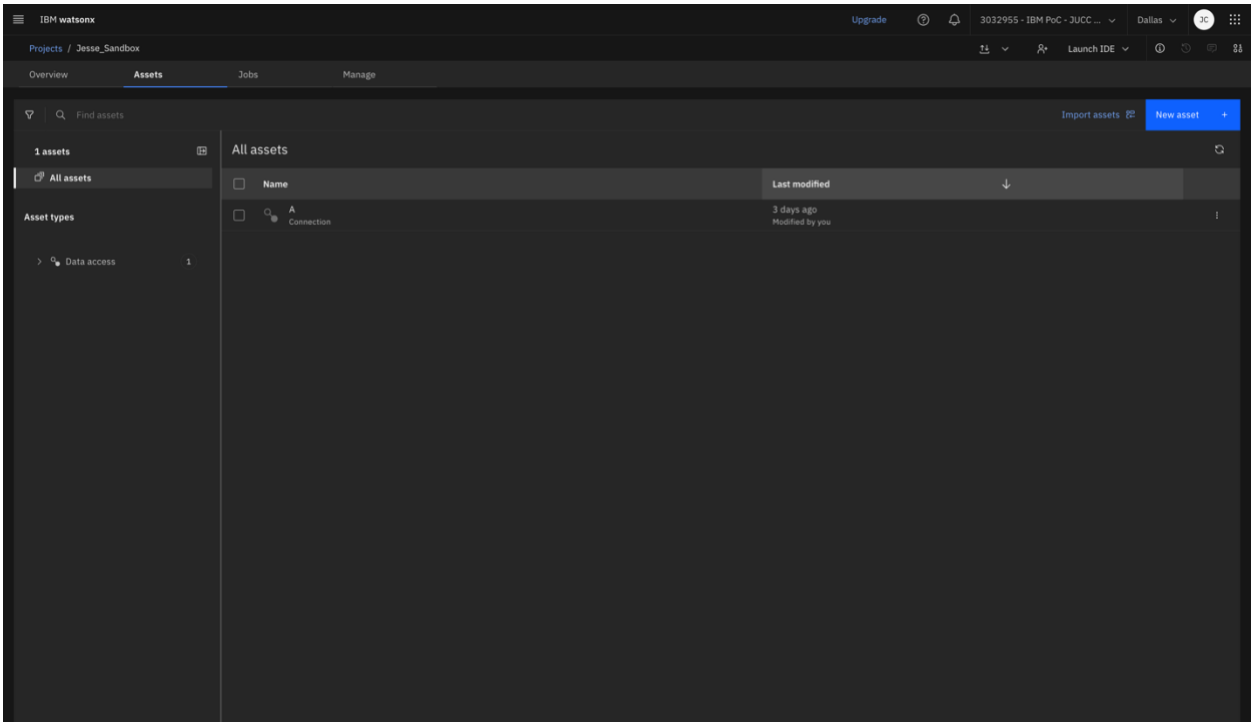

6. In the Import assets dialog, select **Local file** in the left menu.
7. Click the **Data asset** tile.

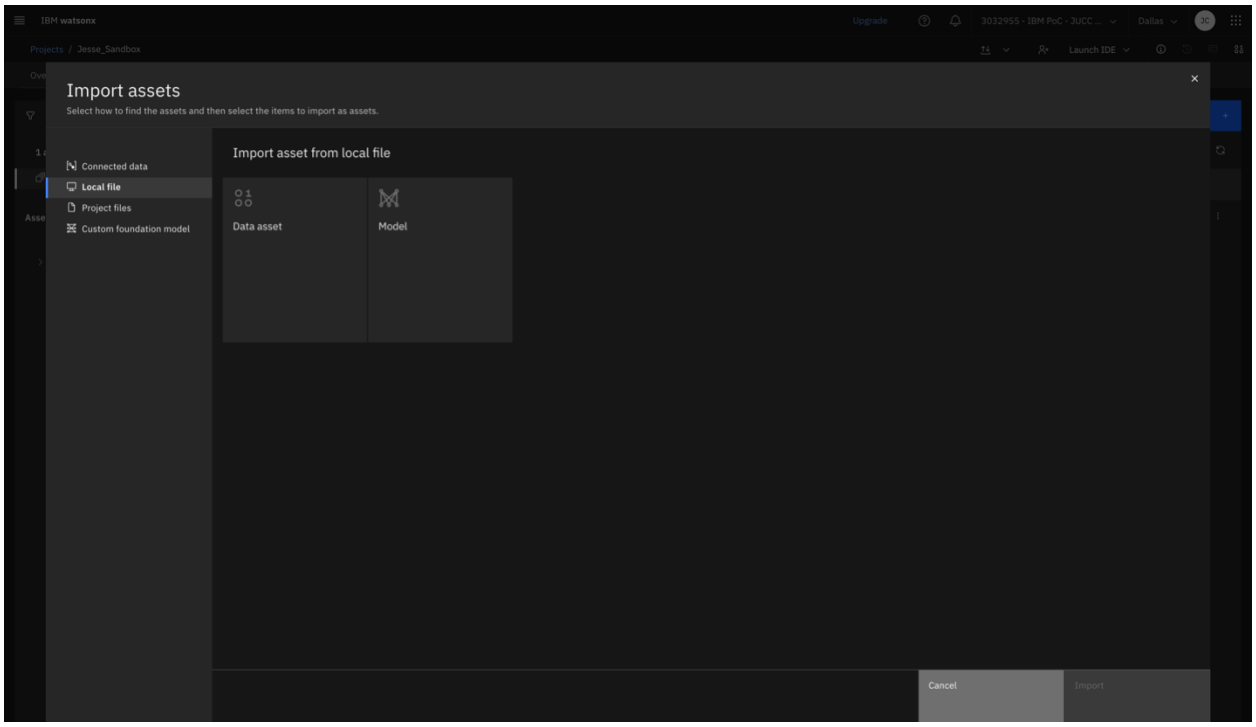

8. Click **Browse**, navigate to and select your dataset file (e.g. my_data.csv).
9. Click **Import** to upload the all data file from "**evaluation_dataset**" folder (Step **3**).

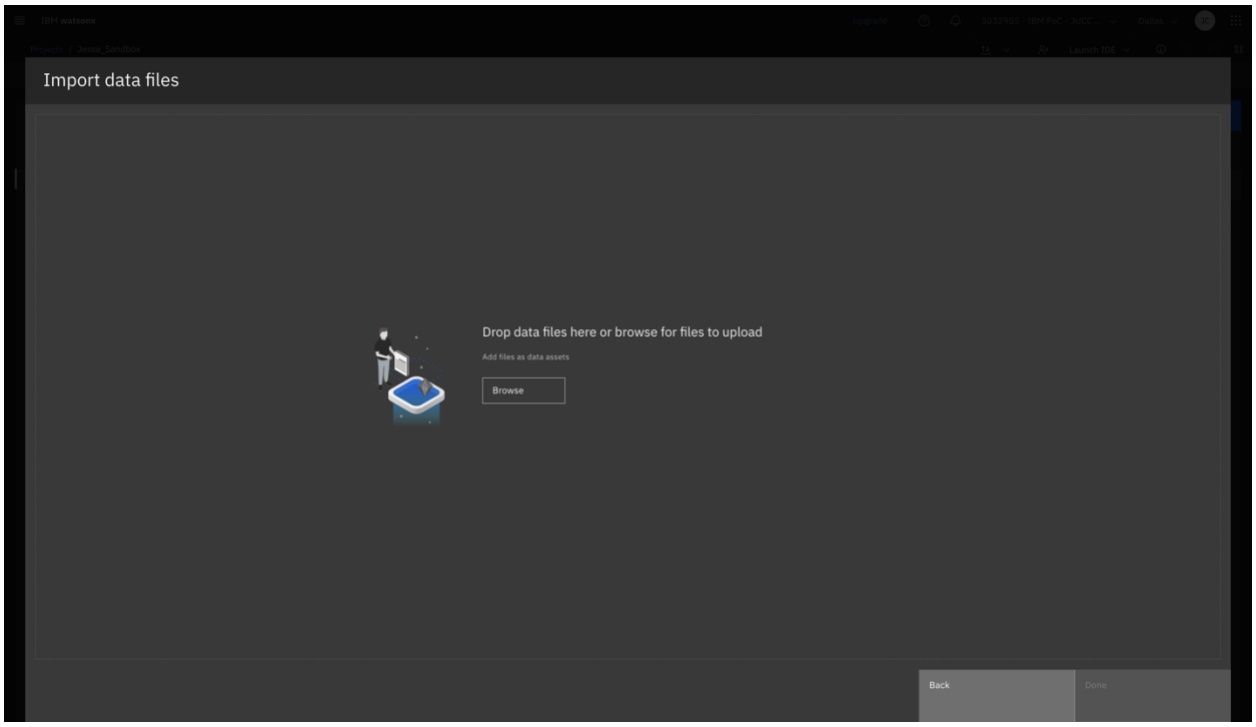

Step 1: Select the Task

1. Navigate to the Assets Page inside projects in IBM watsonx.
2. Select **"Work with models"** from the list of categories.
3. Click on **"Build machine learning or RAG solutions automatically"** using the AutoAI tool.

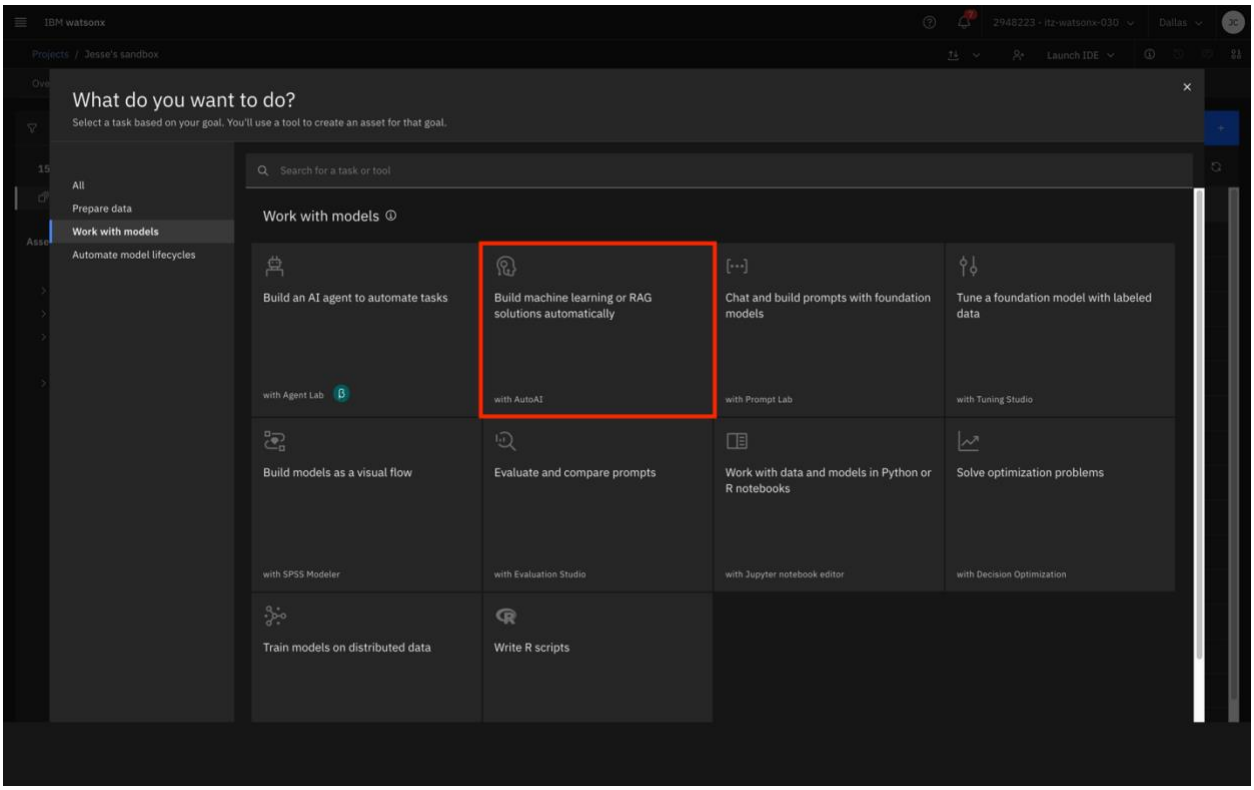

Step 2: Choose the Experiment Type

1. In the **AutoAI** experiment setup screen, choose what you want to build:
    - Select **"Build a RAG solution"** to prepare and optimize RAG patterns based on document collections.
2. Click **Next** to proceed.

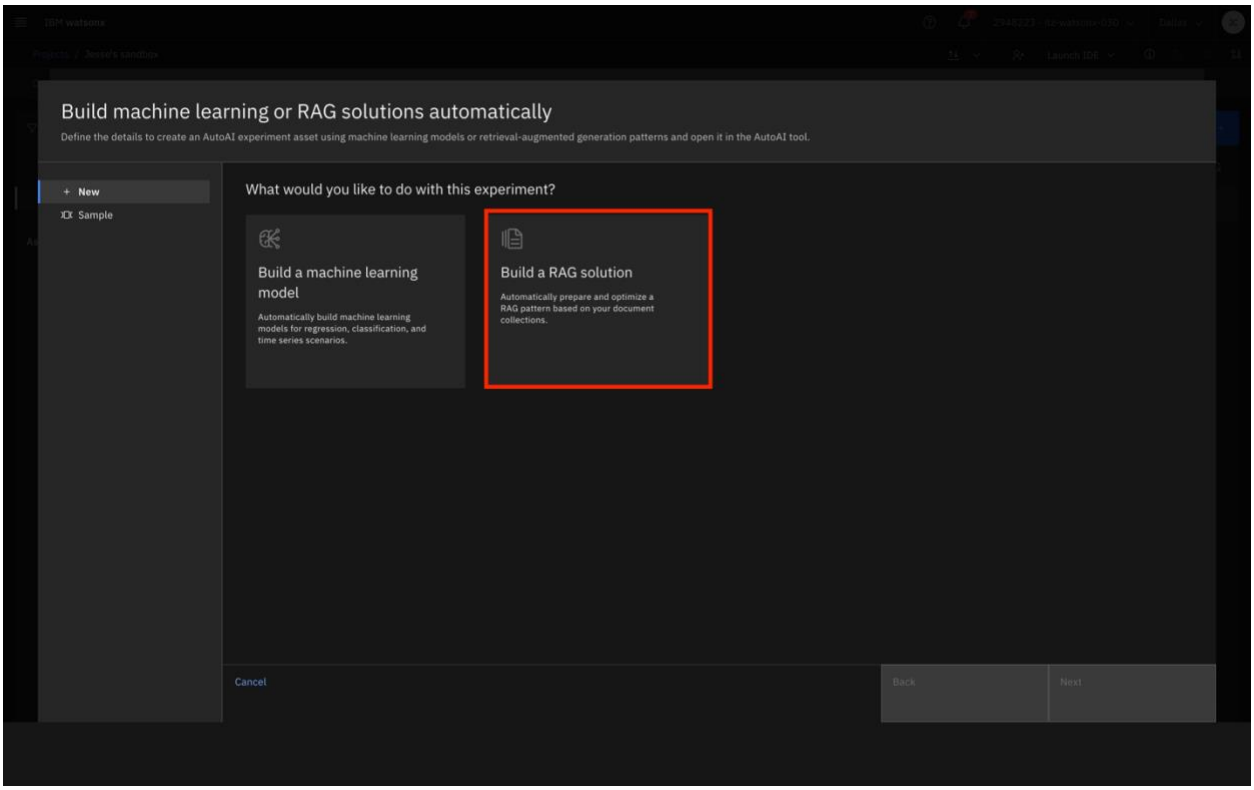

Step 3: Define Experiment Details

1. Provide a name for your experiment (e.g., *AutoAI RAG run - ModelInference documentation*).
2. Optionally, add a description and tags for better organization and searchability.
3. In the configuration section:
    - Select the **watsonx.ai Runtime** service instance.
    - Choose the default environment definition
4. Click **Create** to proceed to the next step.

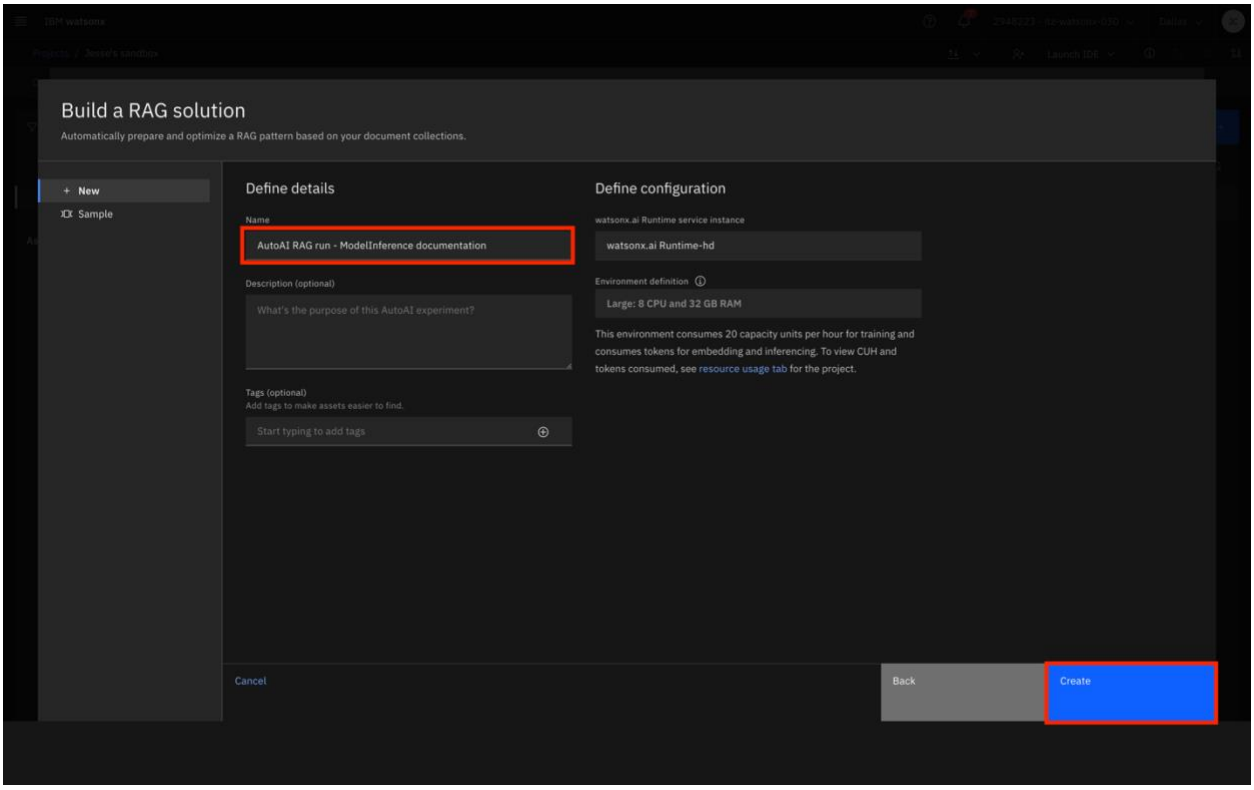

Step 4: Add Document Sources

1. In the **Add sources** screen, upload or select the required documents for the RAG solution:
    - Click **Select from project**, choose **Data asset** in the left-hand Categories pane to upload files from your project.
2. Select the upload all files and click **Open** to upload it.

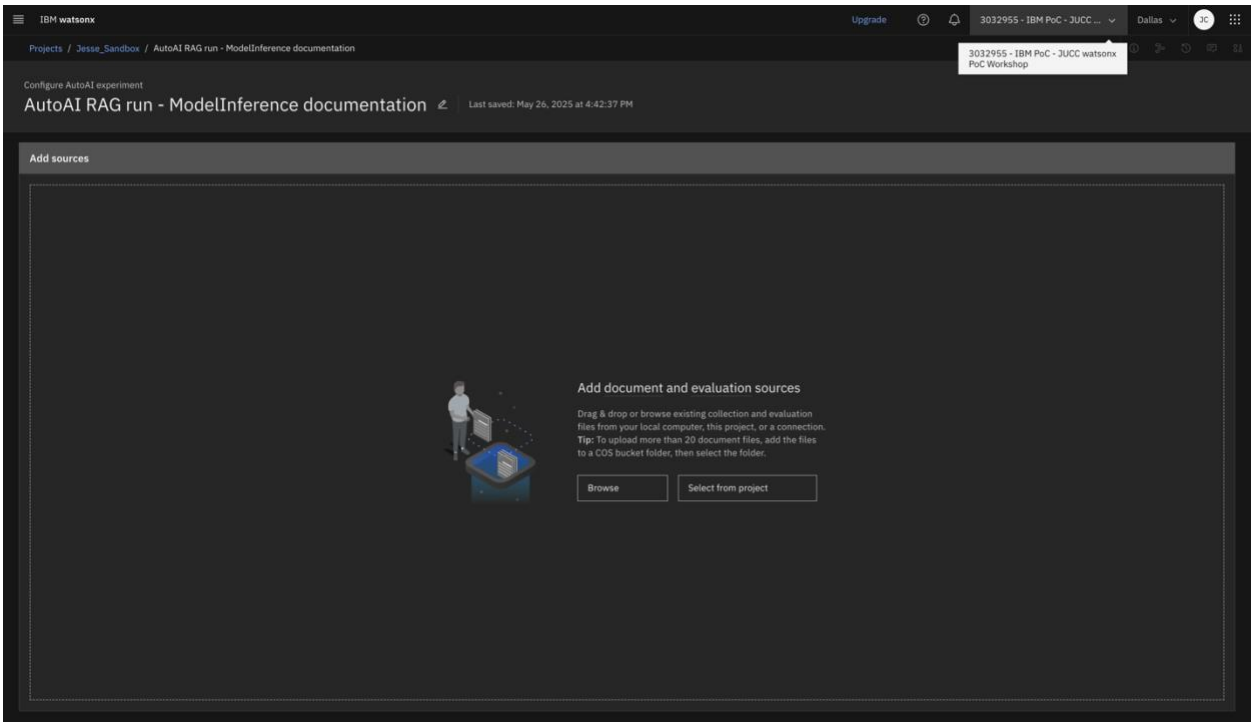

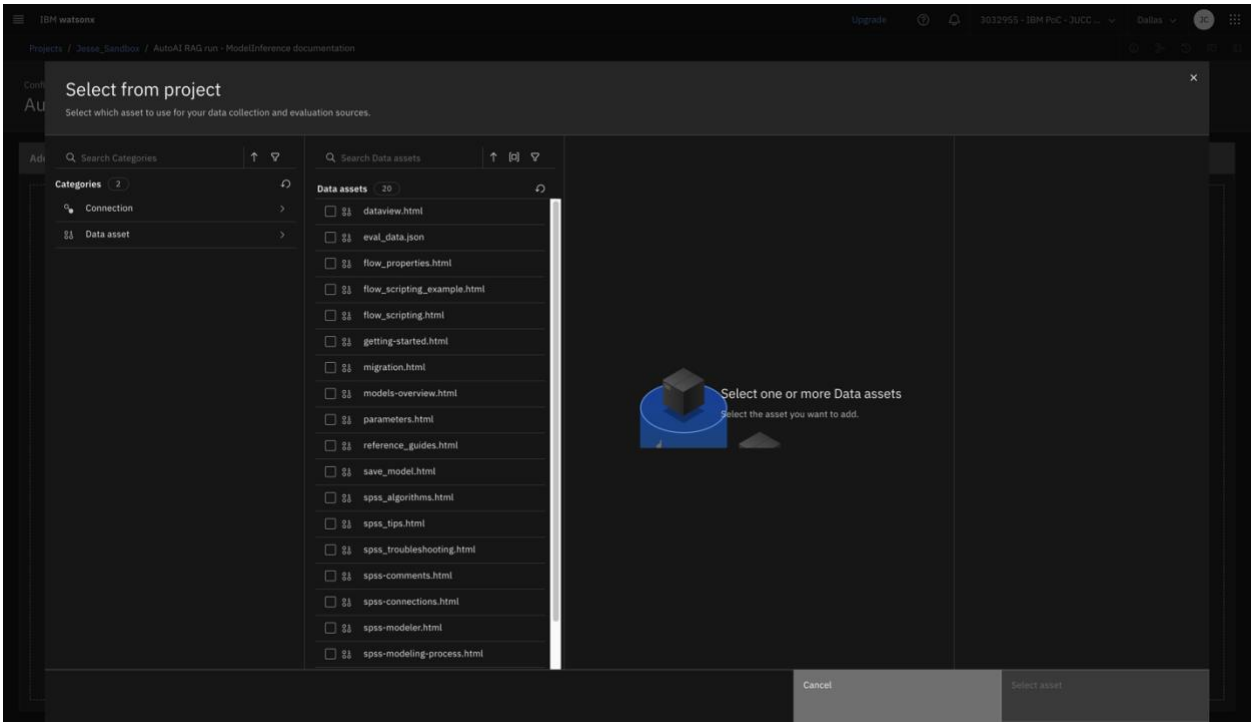
IBM watsonx
Upgrade
3032955 - IBM PoC - JUCC ...
Dallas
JC
Projects / Jesse_Sandbox / AutoAI RAG run - ModelInference documentation
Select from project
Select which asset to use for your data collection and evaluation sources.
Search Categories
Categories 2
Connection
Data asset
Search Data assets
Data assets 20
dataview.html
eval_data.json
flow_properties.html
flow_scripting_example.html
flow_scripting.html
getting-started.html
migration.html
models-overview.html
parameters.html
reference_guides.html
save_model.html
spss_algorithms.html
spss_tips.html
spss_troubleshooting.html
spss-comments.html
spss-connections.html
spss-modeler.html
spss-modeling-process.html
Select one or more Data assets
Select the asset you want to add.
Cancel
Select asset

Step 5: Configure Indexing and Evaluation Sources

1. In the **Configure details** page:
    - For indexing documents:
        - Choose the **'Elasticsearch connection'** to store the vector index.
2. Once the source is uploaded, confirm its inclusion in the experiment.

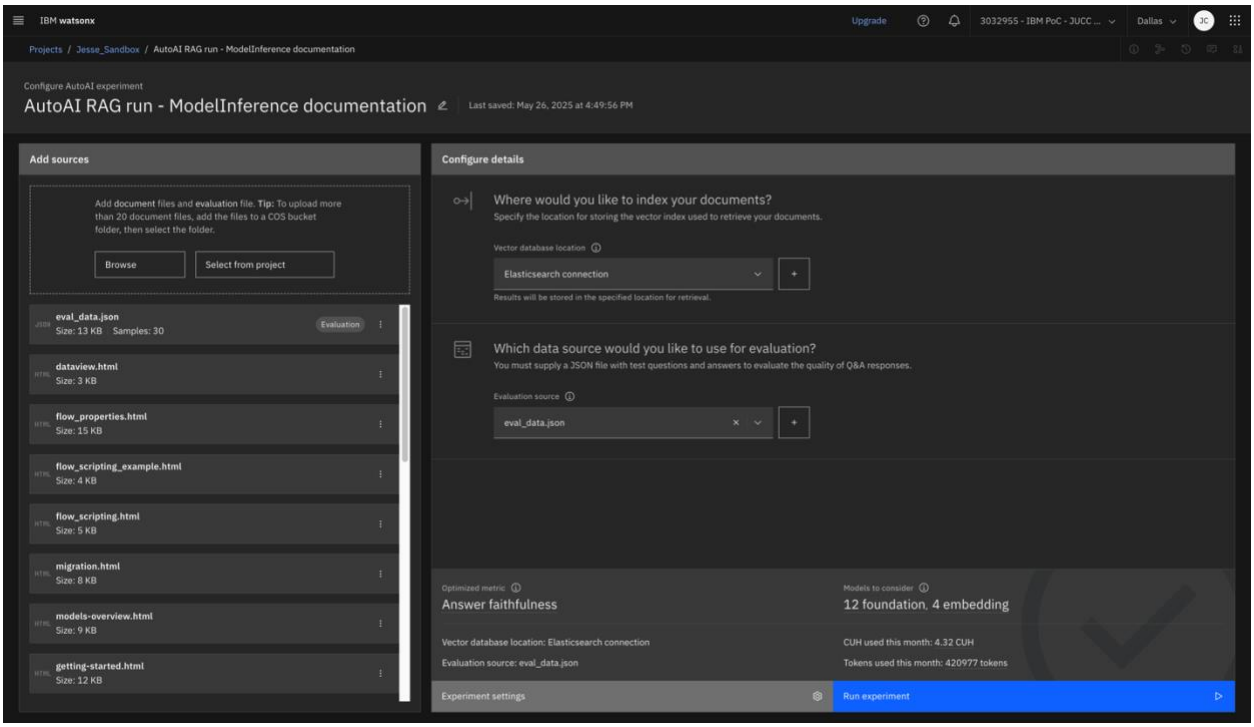

Step 6: Configure Indexing and Experiment Settings

1. Ensure that the documents and evaluation sources are listed correctly in the **Configure details** page.
2. Click **Experiment settings**:

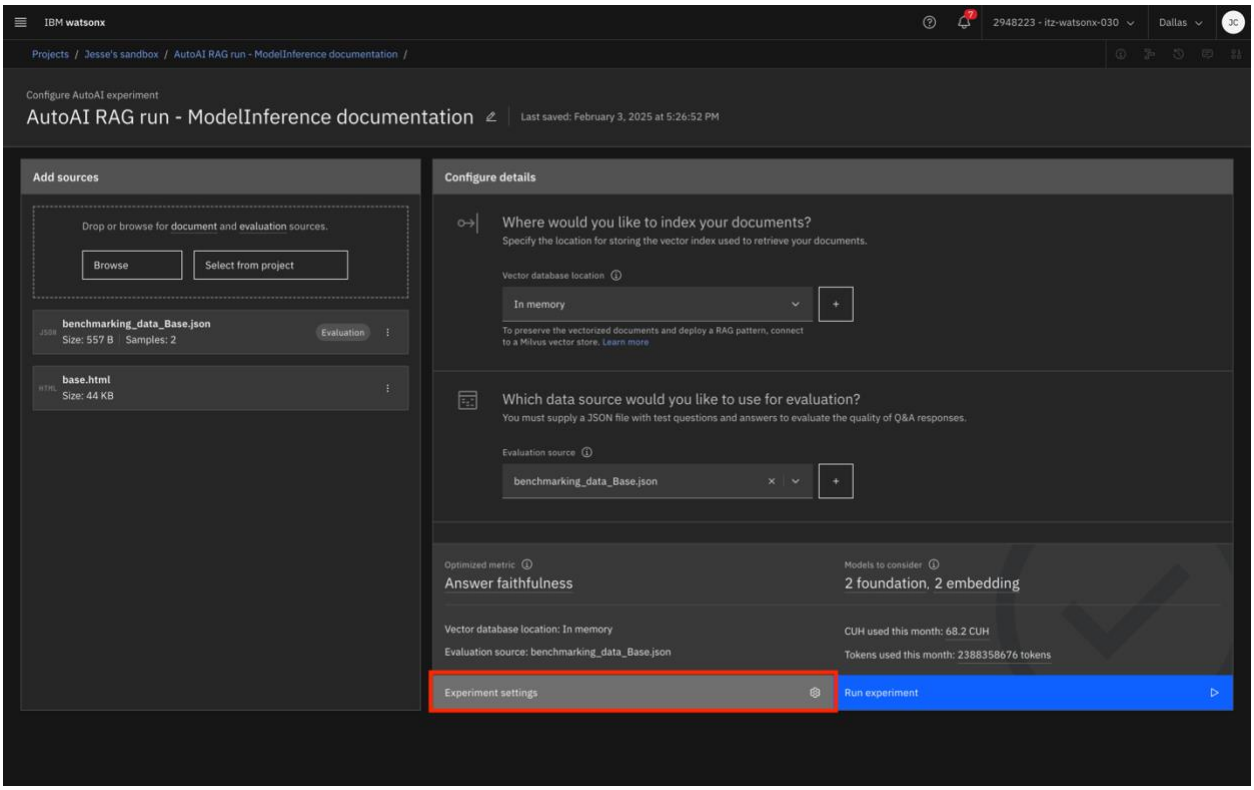

In **Experiment settings**:

- Review the **Optimized metric** (**Answer faithfulness**) to prioritize the quality of generated answers.
- Review the **Retrieval method** (**Window** or **Simple**) for data retrieval.
- Review the **Foundation models** (llama-2-70b-instruct, flan-ul2-20b) for experimentation.

2. Save the changes.

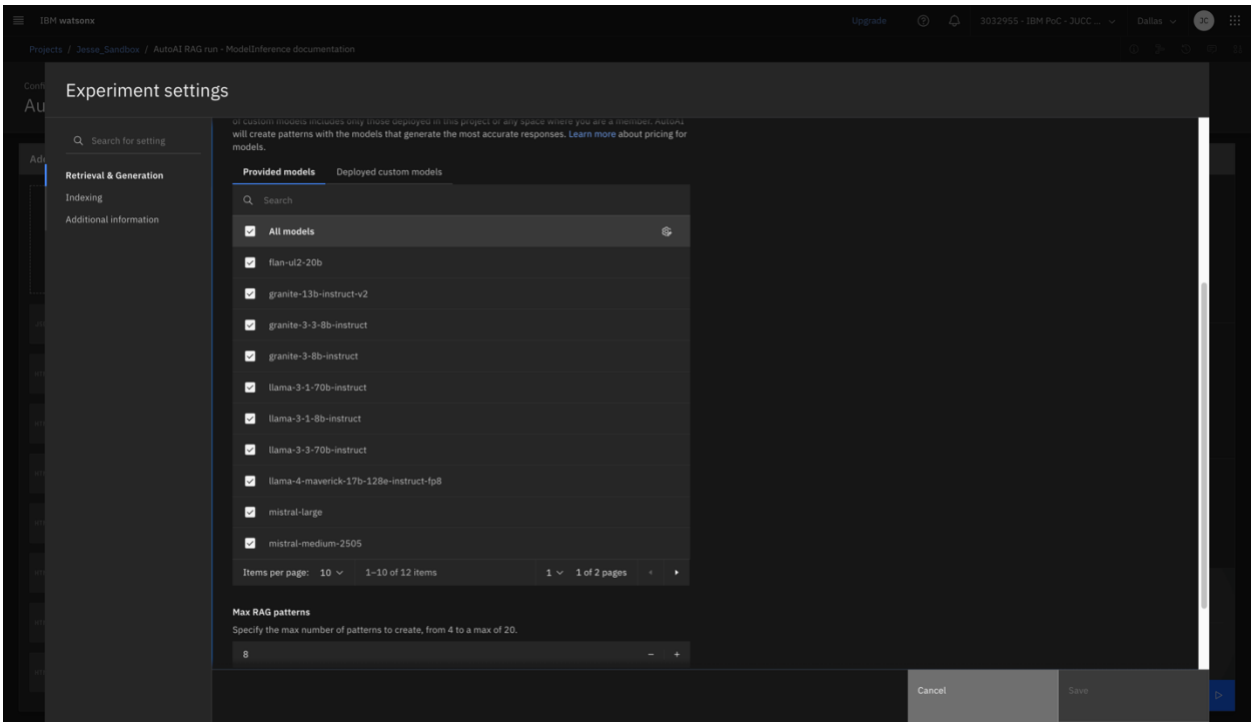

Step 7: Run the Experiment

1. Return to the **Configure details** page.
2. Verify all configurations and sources, then click **Run experiment**.
3. Monitor the execution through the **Progress map**.

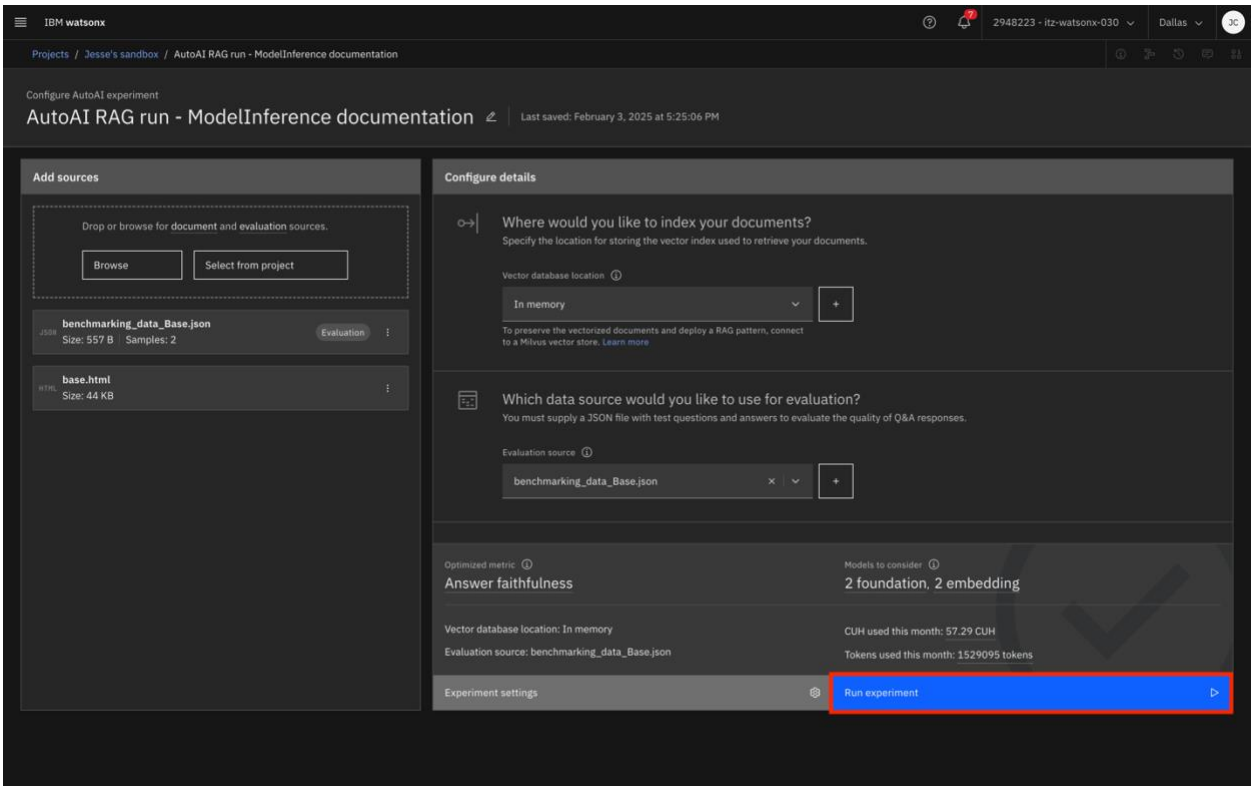

Step 8: Review Results

1. Navigate to the **Experiment summary**:
    - View the **Progress map** to confirm the experiment stages were completed successfully.
    - Check the **RAG Pattern leaderboard** for ranked results based on performance metrics such as **Answer faithfulness** and **Chunk size**.
2. Identify the top-performing pattern, such as Pattern 7 with llama-2-70b-instruct.

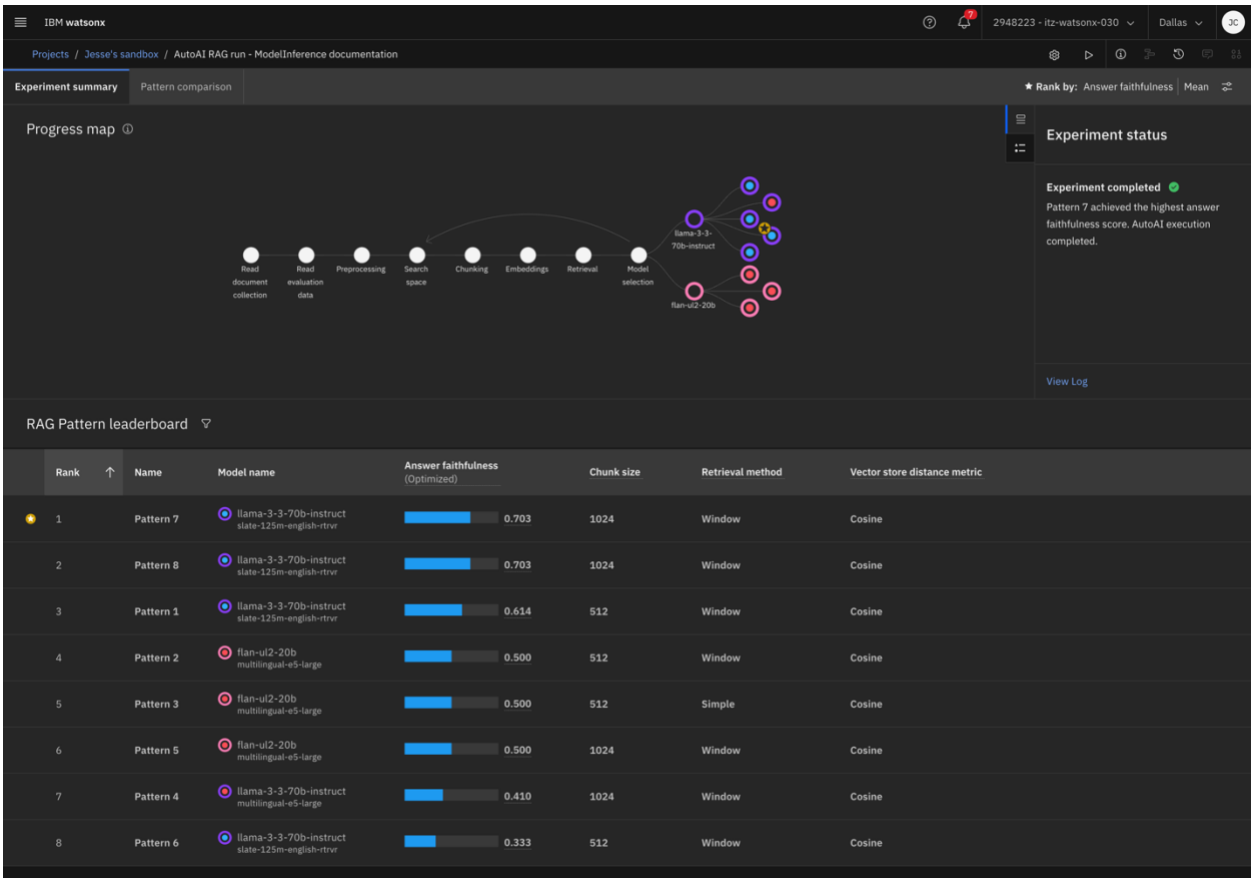

Step 9: Pattern Comparison

1. Switch to the **Pattern comparison** tab for detailed insights:
    - Use the **Metric chart** to compare patterns based on metrics like **Answer faithfulness** and **Correctness**.
    - Analyze the trends to validate the optimal configuration for deployment.

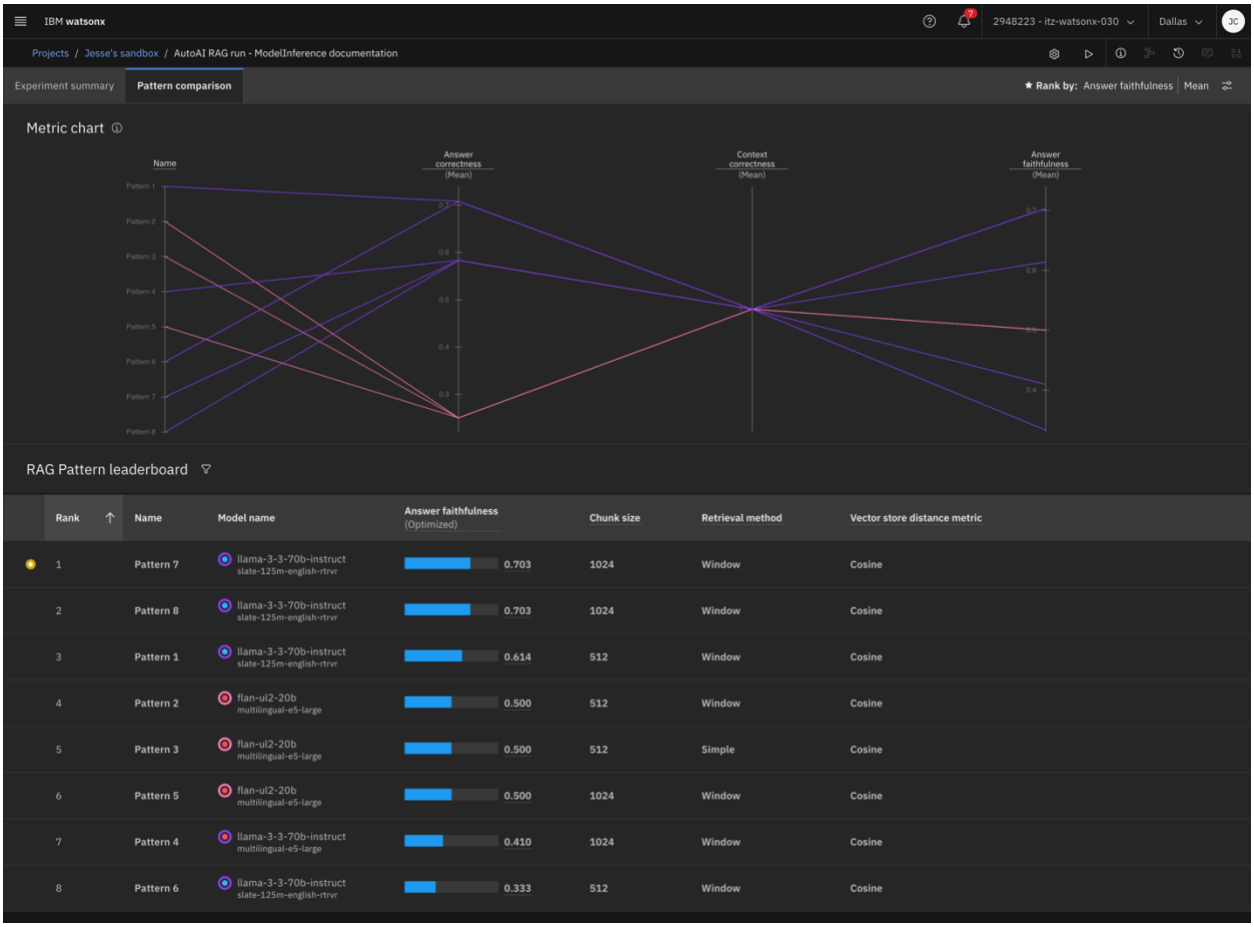

Step 10: Save the Results as a Notebook

1. After the experiment is completed:
    - Review the **RAG Pattern leaderboard** to confirm the top-performing pattern.
2. Click on the **Save as** button to generate a detailed notebook of the experiment results.

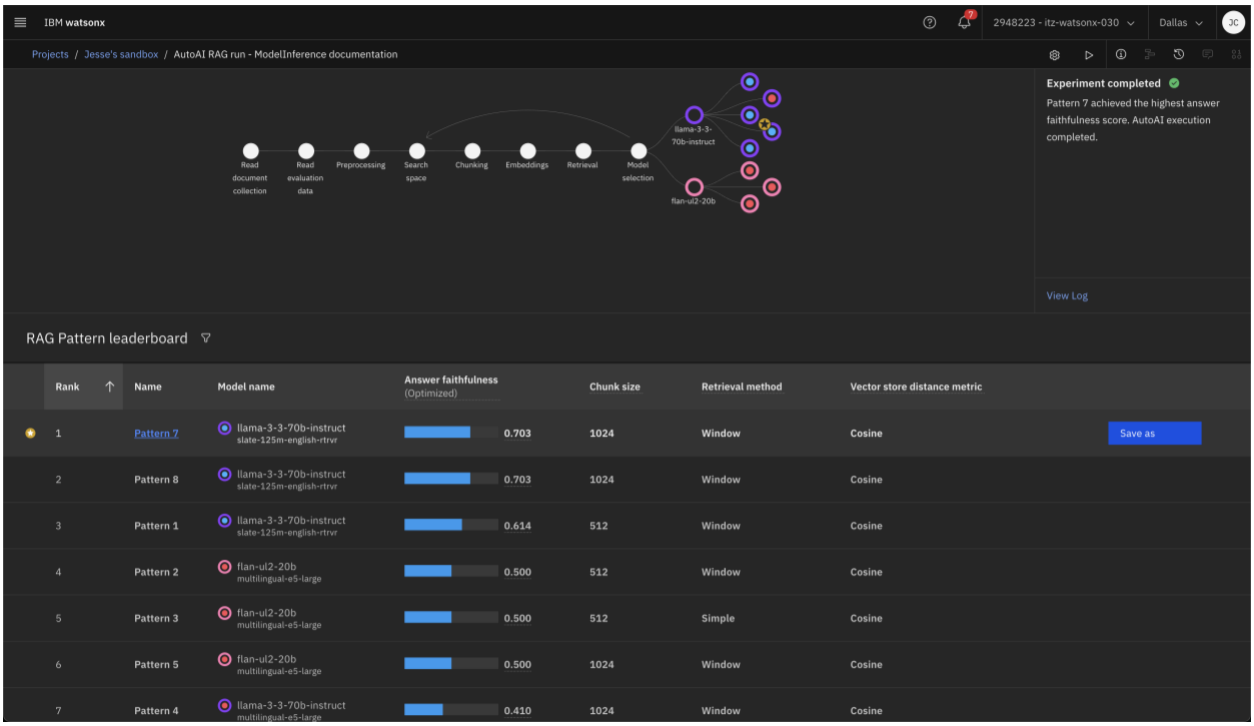

3. In the **Save notebook** dialog:
    - Under Choose your objective, select **Index building** (and Retrieval and generation if you want both).
    - Provide a name for the notebook (e.g., *AutoAI RAG - Top-performing Pattern 7*).
    - Choose the runtime environment and add an optional description or tags for easier identification.
4. Click **Create** to save the notebook for further use or sharing.

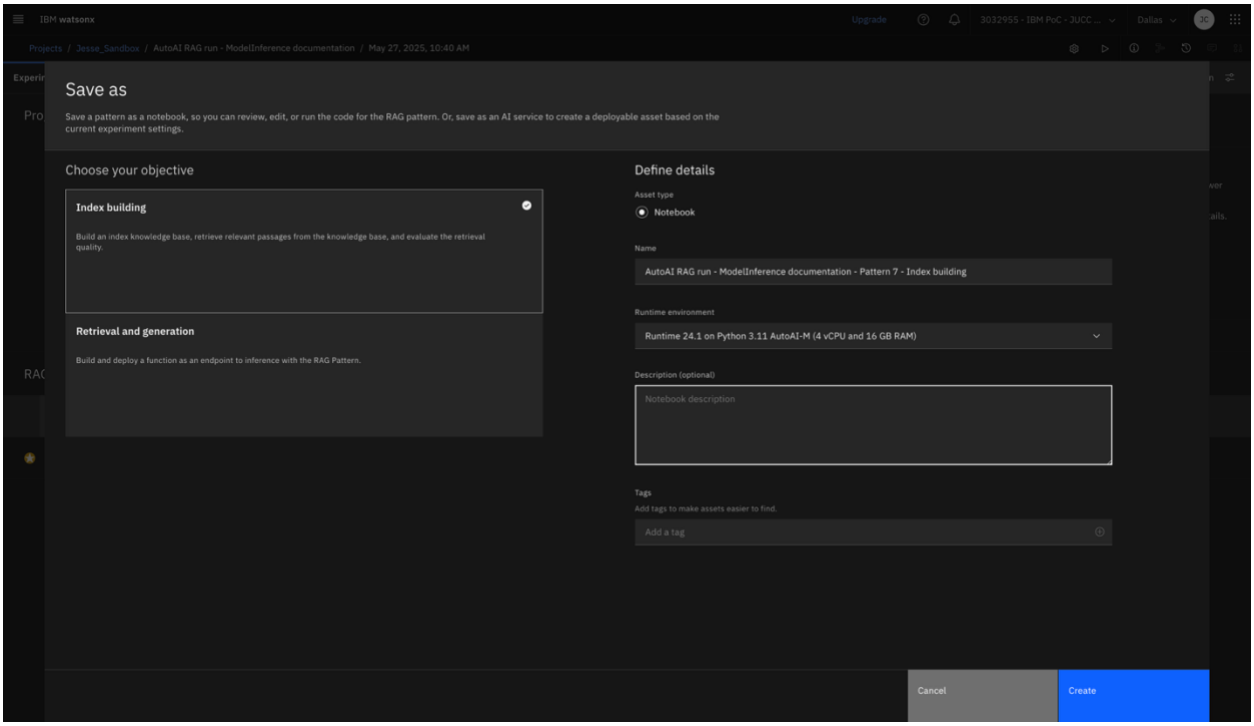

Step 11: Open and Explore the Notebook

1. Once the notebook is created:
    - Open the notebook to explore the detailed steps, data set, configurations, and metrics used in the experiment.

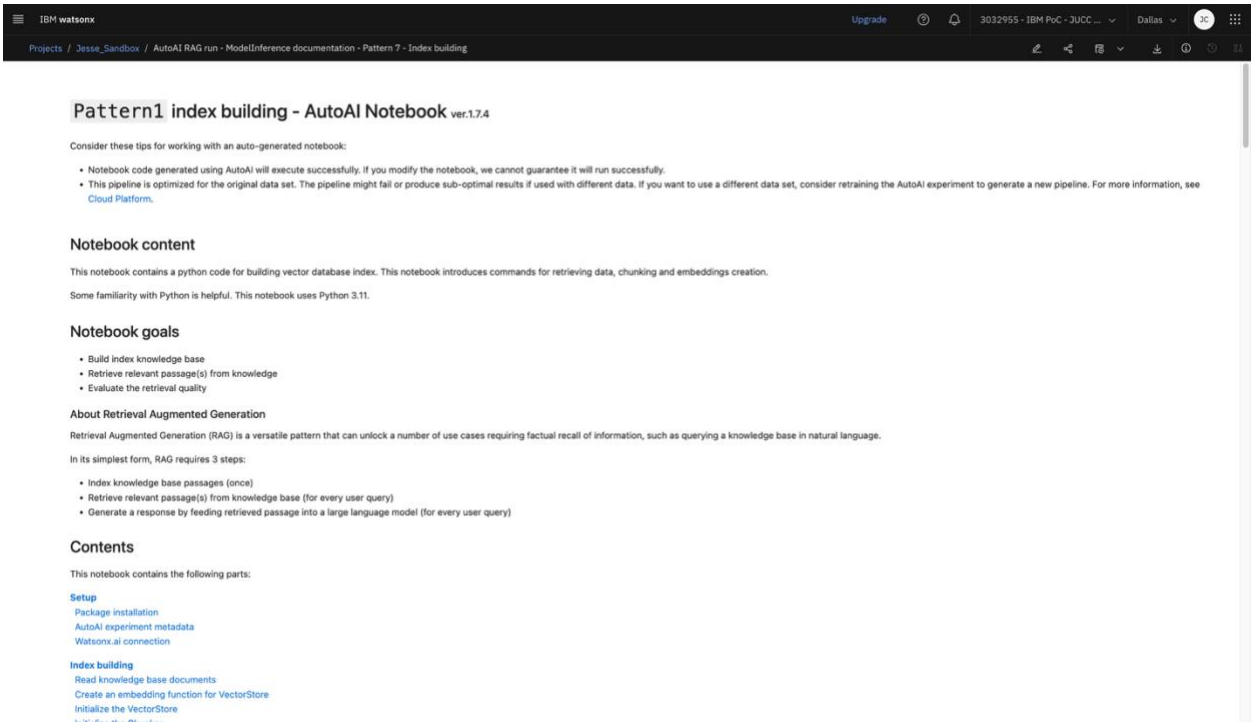

2. Run the notebook with following:
    - **Run the notebook cells**, check Instructions for setting up the vector store and index building.
    - **Watsonx.ai API key** for connection **(step 0)**
    - Details on RAG retrieval and generation processes.
    - Code snippets to reproduce the experiment and generate results.
3. Use the notebook to further refine or replicate the experiment as needed.

Step 12: Extract RAG Pattern Configuration Keys

1. Open the code or notebook cell containing your rag_pattern = { ... } definition.

2. Locate the top-level keys and the nested settings that drive each module.
3. Copy out the following key names for use in your module in later lab:

IBM watsonx.ai Embeddings Key

	Where to get it?	Value
watsonx API Endpoint	->	https://us-south.ml.cloud.ibm.com
API Key	Step 0	(get your own config)
Project ID	Find 'experiment_metadata'	(get your own config) e.g. cd2dadd2-68e5-4680-8d79-cf8db08cc8c8
Embedding Model Name	rag_pattern: embeddings -> model_id	(get your own config) e.g. intfloat/multilingual-e5-large

IBM watsonx.ai Key

	Where to get it?	Value
watsonx API Endpoint	->	https://us-south.ml.cloud.ibm.com
Model Name	rag_pattern: generation -> model_id	(get your own config) e.g. meta-llama/llama-3-3-70b-instruct

Elasticsearch Vector Store

	Where to get it?	Value
Elasticsearch URL	->	https://us-south.ml.cloud.ibm.com
Index Name	rag_pattern: vector_store -> index_name	(get your own config) e.g. autoai_rag_69210801_20250527023933
Username	Step 0 -> sub-step 6	(get your own config) e.g. ibm_cloud_c06543e4_08e4_42d9_836e_9de74a8cff8d
Password	Step 0 -> sub-step 6	(get your own config) e.g. eGq3GJdEh9kHhQHfetUu3gOE6hwGIP7l

Step 13: Upload the ES Flow into Lang flow

1. In your browser, navigate to the Langflow UI: http://158.176.6.193:7860/
2. In the left-hand sidebar, click Upload a flow.
3. browse to the local JUCC-Workshop folder you downloaded earlier, select **ES Flow Without Keys.json**, and click Open.
4. Langflow will import the flow and display it on the canvas for you to review and configure

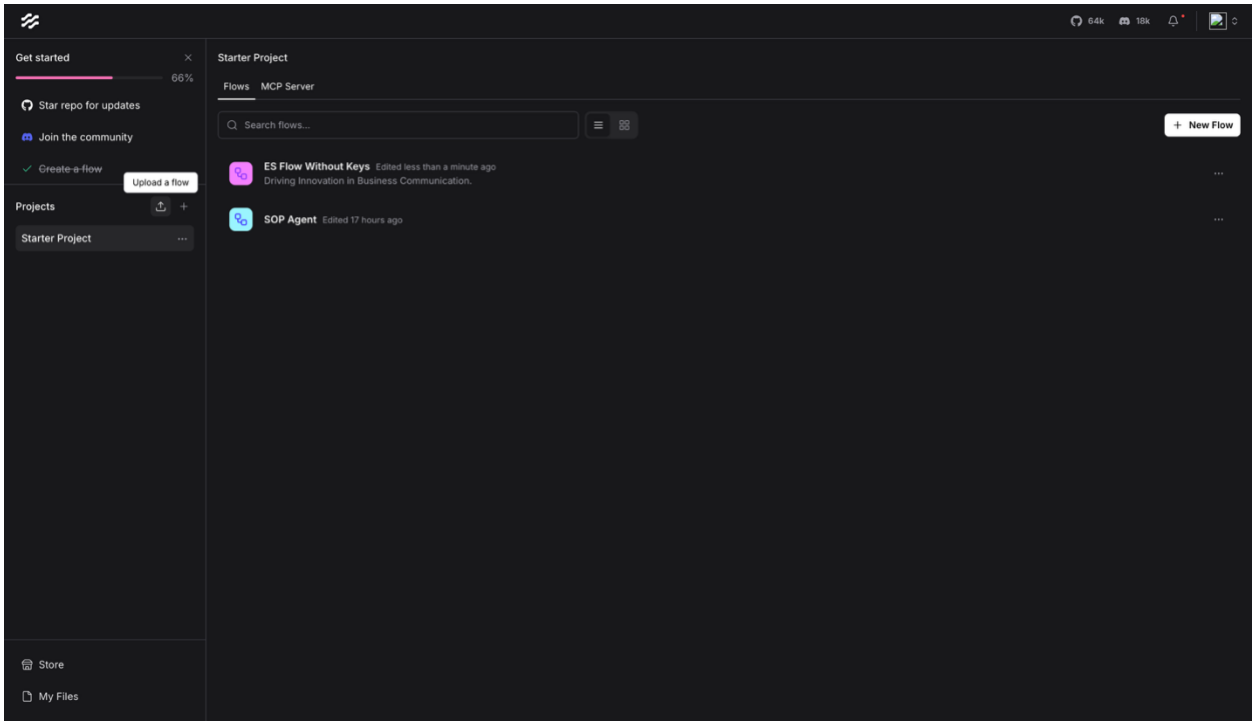

Step 14: Configure Each Langflow Component and Understand Its Role

1. **watsonx.ai Embeddings**
   Role: Transforms text (documents or queries) into vector representations that capture semantic meaning, enabling similarity-based search in later steps.
   - watsonx API Endpoint: https://us-south.ml.cloud.ibm.com
   - watsonx project id: your project ID
   - API Key: your API key
   - Model Name: your embeddings model (e.g. intfloat/multilingual-e5-large)

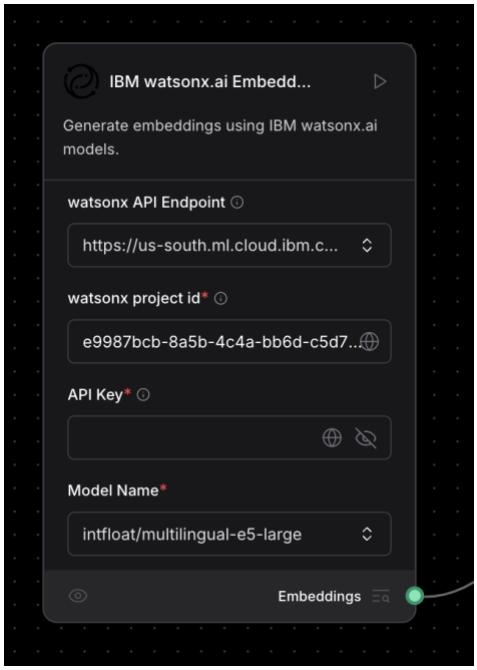

2. Elasticsearch Vector Store
   Role: Stores and indexes those embeddings; given a query vector, it retrieves the top most similar document chunks to serve as context.
   - Elasticsearch URL: your Elasticsearch URL
   - Index Name: your index name (e.g. autoai_rag_69210801_20250527023933)
   - Username: your Elasticsearch username
   - Password: your Elasticsearch password

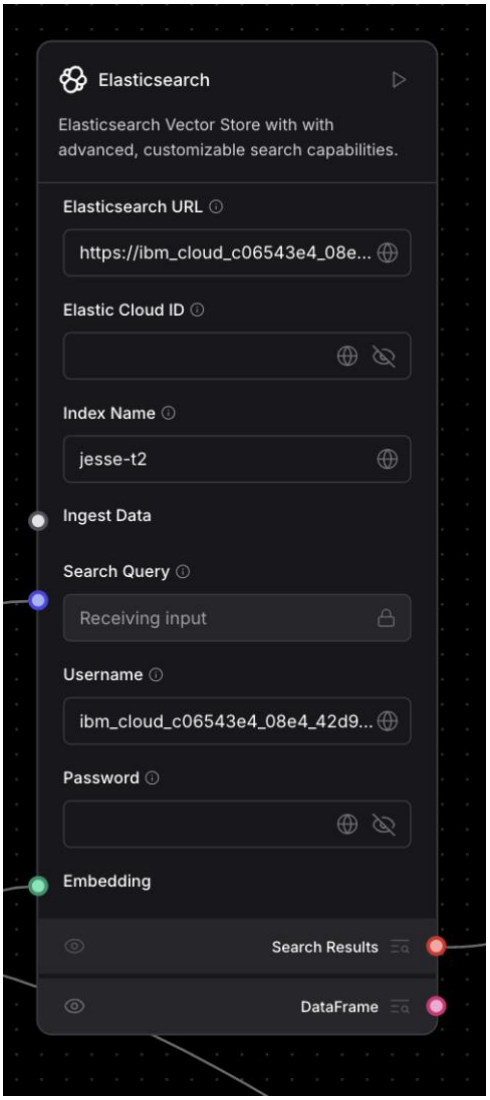

3. IBM watsonx.ai (Generation)
   Role: Consumes the retrieved chunks plus the user's query to generate a coherent, context-grounded answer using a foundation language model.
   - watsonx API Endpoint: https://us-south.ml.cloud.ibm.com
   - Model Name: your generation model (e.g. meta-llama/llama-3-3-70b-instruct)

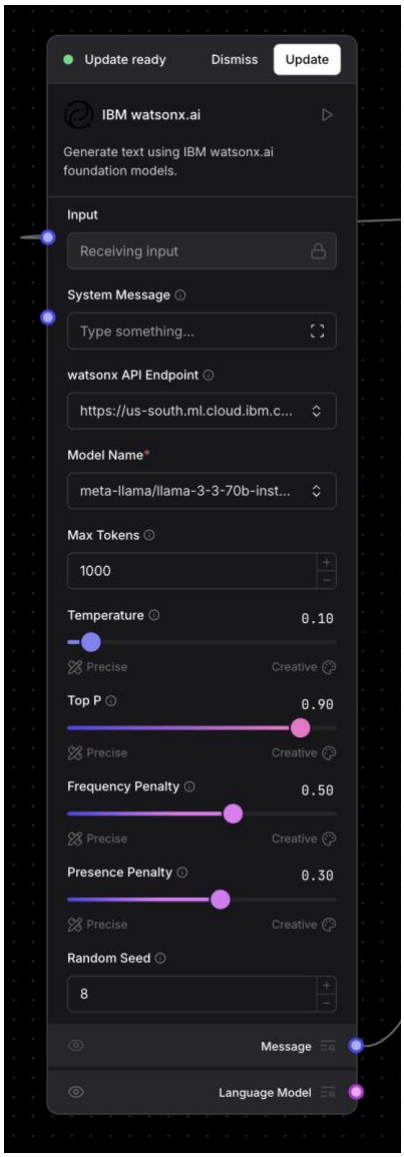

Step 15: Test Your RAG Flow in the Playground

1. In the top-right corner of Langflow, click **Playground**.
2. In the "New chat" prompt that appears, type your test question (for example, "What is SPSS?").
3. Press **Enter** or click **Send**.
4. Observe how the flow:
    - Converts your question into an embedding,
    - Retrieves the most relevant document chunks from Elasticsearch,
    - Feeds those chunks + your question into the generation model,
    - Returns a grounded answer in the chat pane.
5. Review the **Session** sidebar to inspect earlier queries and responses or start a fresh chat by clicking the + next to **Chat**.

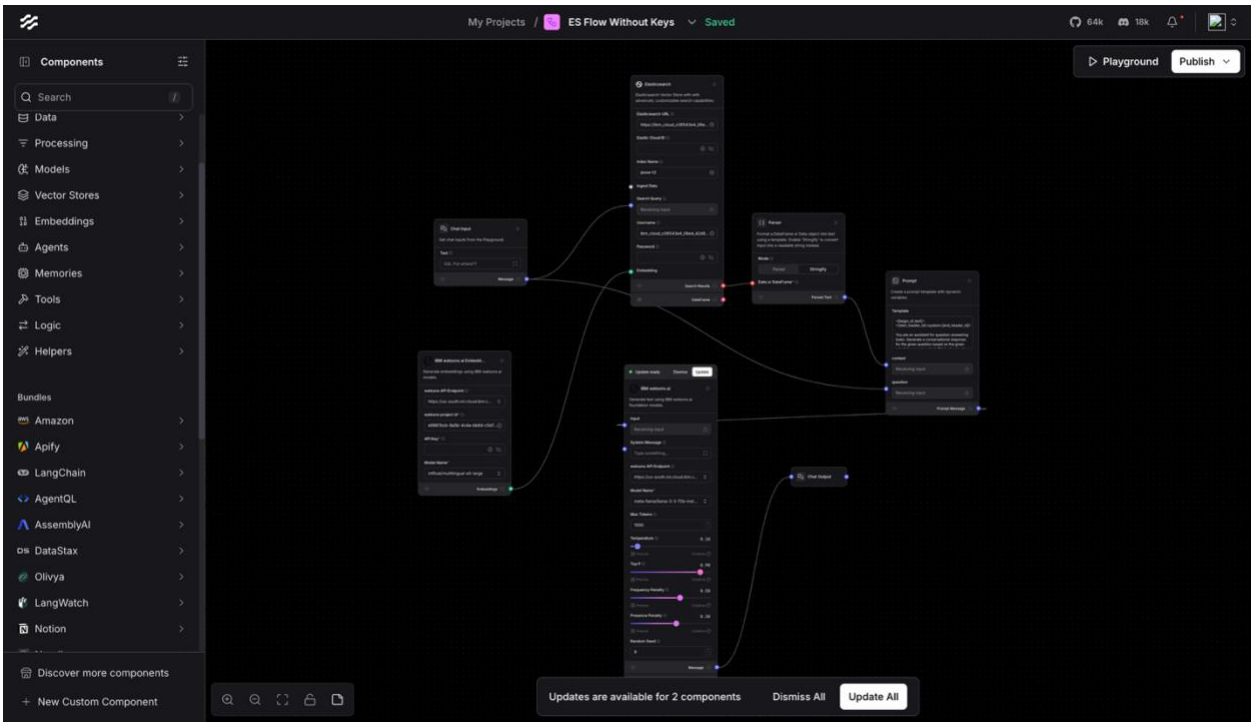

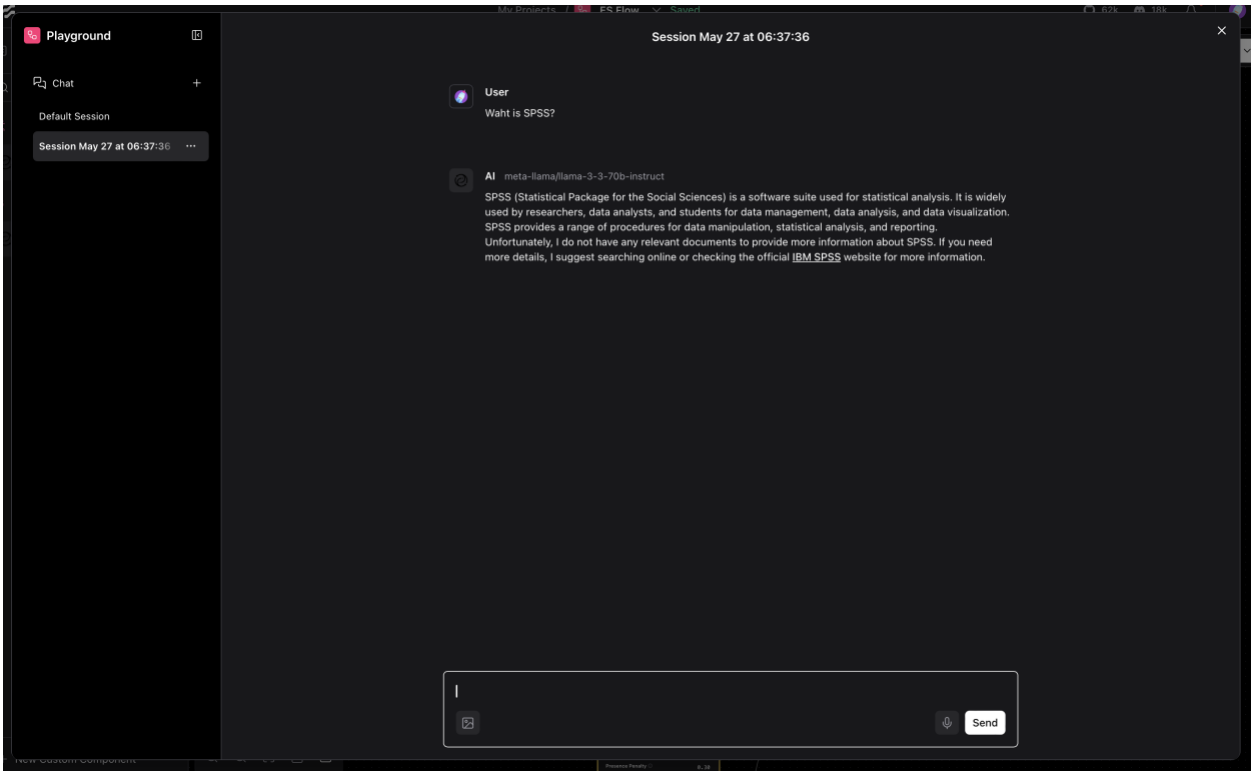
Playground
Chat
Default Session
Session May 27 at 06:37:36
Session May 27 at 06:37:36
User
Waht is SPSS?
AI
meta-llama/llama-3-3-70b-instruct
SPSS (Statistical Package for the Social Sciences) is a software suite used for statistical analysis. It is widely used by researchers, data analysts, and students for data management, data analysis, and data visualization. SPSS provides a range of procedures for data manipulation, statistical analysis, and reporting.
Unfortunately, I do not have any relevant documents to provide more information about SPSS. If you need more details, I suggest searching online or checking the official IBM SPSS website for more information.
Send